

Assignment Eight

Sybil Melton

November 13, 2014

CONTENTS

1	MovieLens	3
1.1	Highest Average	3
1.1.1	Solution	3
1.1.2	Result	4
1.2	Most Ratings	5
1.2.1	Solution	5
1.2.2	Result	5
1.3	Highest Average by Women	6
1.3.1	Solution	6
1.3.2	Result	7
1.4	Highest Average by Men	7
1.4.1	Solution	8
1.4.2	Result	8
1.5	Top Gun Correlation	8
1.5.1	Solution	8
1.5.2	Result	10
1.6	Top Raters	10
1.6.1	Solution	10
1.6.2	Result	11
1.7	User Agreement	11
1.7.1	Solution	11
1.7.2	Result	12
1.8	User Disagreement	12
1.8.1	Solution	12
1.8.2	Result	12
1.9	Men by Age	13
1.9.1	Solution	13
1.9.2	Result	14
1.10	Women by Age	16
1.10.1	Solution	16
1.10.2	Result	16
1.11	Python Code	18

LISTINGS

1	Get Highest Average Modules	3
2	Get Most Ratings Code	5
3	Get Top Ratings by Gender Module	6
4	Top Gun Correlation Module	9
5	Top Raters Module	10
	recommendations.py	11
6	User Disagreement	12
7	Ratings by Age	13
8	Python Recommendations	18

1 MOVIELENS

The goal of this project is to use the basic recommendation principles we have learned for user-collected data. You will modify the code given to you which performs movie recommendations from the MovieLense data sets.

The MovieLense data sets were collected by the GroupLens Research Project at the University of Minnesota during the seven-month period from September 19th, 1997 through April 22nd, 1998. It is available for download from <http://www.grouplens.org/node/73>

There are three files which we will use:

1. u.data
2. u.item
3. u.user

The code for reading from the u.data and u.item files and creating recommendations is described in the book Programming Collective Intelligence (check email for more details). You are to modify recommendations.py to answer the following questions. Each question your program answers correctly will award you 10 points. You must have the question answered completely correct; partial credit will only be awarded if your answer is very close to the correct one.

1.1 HIGHEST AVERAGE

What 5 movies have the highest average ratings? Show the movies and their ratings sorted by their average ratings.

1.1.1 SOLUTION

The first step was to load the data into a usable structure. The movie title to id mappings were used frequently, so getMovieTitles was written from the loadMovieLens function, to be reusable and returns a dictionary of the mappings. Module loadMovieRatings is a modified version of loadMovieLens. It was written to take the user ratings from u.data, combine it with the movie titles, and return a dictionary of the movie to user rating mappings with the 'ratings' key. [4] This key was used to keep track of the averages later on. Since there were several requirements for rating averages, ratingAverage was developed for calculating and returning the average of a list of ratings, based on the 'ratings' key. There were also several instances of printing top rating averages, so printRatings was written to print each result. The To cover the event that there was a tie of more than five, getTop allowed for printing all of the items that have the top result. If that list was less than five results, it continued to look to fill the requirement. The module getRatings combined all of these functions. The movie to user rating mappings were loaded into a dictionary, then for each movie, the average was added to the dictionary with the 'average' key. The dictionary was turned into a list of (key, value) tuples, and sent to printRatings. [1] The formatting used left and right justification to get an easier to read output. [3] Listing 1 is the functions used.

```
1 def getMovieTitles(path='data'):  
2     # Get movie titles  
3     movies = {}  
4     for line in open(path + '/u.item'):  
5         (id, title) = line.split('|')[0:2]  
6         movies[id] = title  
7     return movies  
8  
9 def loadMovieRatings(path='data'): #load movies with title and ratings list  
10    # Get movie titles  
11    movies = getMovieTitles()  
12    # Get all ratings for each movie
```

```

13 rate = {}
14 for line in open(path + '/u.data'):
15     (user, movieid, rating, ts) = line.split('\t')
16     rate.setdefault(movieid, {})
17     rate[movieid]['title'] = movies[movieid]
18     if rate[movieid].has_key('ratings'):
19         rate[movieid]['ratings'].append(float(rating))
20     else:
21         rate[movieid]['ratings'] = []
22         rate[movieid]['ratings'].append(float(rating))
23 return rate
24
25 def ratingAverage(ratings): #get average of a list of movie ratings
26     sum = 0.0
27     for i in range(0, len(ratings['ratings'])):
28         sum = sum + ratings['ratings'][i]
29     avg = sum / len(ratings['ratings'])
30     return avg
31
32 def getTop(sort, type, n=5):
33     max = sort[0][1][type]
34     top = []
35     top.append(sort[0])
36     for s in range(1, len(sort)):
37         if sort[s][1][type] == max:
38             top.append(sort[s])
39         else:
40             break
41     while len(top) < n:
42         max = sort[len(top)][1][type]
43         for s in range(len(top), len(sort)):
44             if sort[s][1][type] == max:
45                 top.append(sort[s])
46             else:
47                 break
48     return top
49
50 def printRatings(items, type, n): #print list of ratings
51     sort = sorted(items, key=lambda x: x[1][type], reverse=True)
52     col_width = 0
53     top = getTop(sort, type, n)
54     # Get longest title for printed column width
55     for x in range(0, len(top)):
56         if len(top[x][1]['title']) > col_width:
57             col_width = len(top[x][1]['title']) + 2
58     # print the result
59     print string.ljust('Title', col_width), string.rjust('Average Rating', 3)
60     print string.ljust('\n-----', col_width), string.rjust('-----', 3)
61     for x in range(0, len(top)):
62         print string.ljust(top[x][1]['title'], col_width), string.rjust(str(top[x][1][type]), 3)
63
64 def getRatings(n=5): #get top average, most rated
65     rate = loadMovieRatings()
66     # Get the average ratings per movie
67     for r in rate:
68         rate[r]['average'] = ratingAverage(rate[r])
69     print 'Top ' + str(n) + ' movies by Average Rating:'
70     printRatings(rate.items(), 'average', n)

```

Listing 1: Get Highest Average Modules

1.1.2 RESULT

The top movies all had a rating of 5.0. In my testing I chose a number higher than 5, to show that my functions were working correctly. The results were saved into results.txt and are included in this report. There were more than five, but for the purposes of the assignment, the output still states “Top 5.” Figure 1.1 shows the list.

Top 5 movies by Average Rating:	
Title	Average Rating
-----	-----
Entertaining Angels: The Dorothy Day Story (1996)	5.0
Star Kid (1997)	5.0
Great Day in Harlem, A (1994)	5.0
They Made Me a Criminal (1939)	5.0
Someone Else's America (1995)	5.0
Saint of Fort Washington, The (1993)	5.0
Aiqing wansui (1994)	5.0
Santa with Muscles (1996)	5.0
Prefontaine (1997)	5.0
Marlene Dietrich: Shadow and Light (1996)	5.0

Figure 1.1: Highest Average Ratings

1.2 MOST RATINGS

What 5 movies received the most ratings? Show the movies and the number of ratings sorted by number of ratings.

1.2.1 SOLUTION

The same function, `getRatings`, was used as in the first question. The data was sorted differently, on the length of the ratings list, for each movie. Since the sorting had to be completed by the rating list length, instead of an average value, the print statements had to be written separately. Listing 2 is the additional code used in the `getRatings` function.

```

1  # Get the most rated movies
2  rsort = sorted(rate.items(), key=lambda x: len(x[1]['ratings']), reverse=True)
3  print '\nTop ' + str(n) + ' movies by most ratings:'
4  col_width = 0
5  for x in range(0, n):
6      if len(rsort[x][1]['title']) > col_width:
7          col_width = len(rsort[x][1]['title']) + 2
8  print string.ljust('Title', col_width), string.rjust('Number of Ratings', 3)
9  print string.ljust('-----', col_width), string.rjust('-----', 3)
10 for x in range(0, n):
11     print string.ljust(rsort[x][1]['title'], col_width), string.rjust(str(len(rsort[x][1]['ratings'])),
3)

```

Listing 2: Get Most Ratings Code

1.2.2 RESULT

There were no surprises to me, in this list of top rated movies, since I have seen all of them. I did not see any ties, so for the assignment, it was not included. Figure 1.2 shows the output.

Top 5 movies by most ratings:

Title	Number of Ratings
Star Wars (1977)	583
Contact (1997)	509
Fargo (1996)	508
Return of the Jedi (1983)	507
Liar Liar (1997)	485

Figure 1.2: Most Ratings

1.3 HIGHEST AVERAGE BY WOMEN

What 5 movies were rated the highest on average by women? Show the movies and their ratings sorted by ratings.

1.3.1 SOLUTION

A separate module was written for the gender classifications, because additional user information was needed. Module `loadUserInfo` placed the user id to movie rating mappings with a nested dictionary with age and gender. [1] The `loadMovieRatings` module was used again, for consistency and reuse of average and printing functions. A dictionary was created for women, initialized with each movie id. If a user's gender was F, for female, the movie list was checked if the user had rated it and the rating was added to the women dictionary. Once all the women ratings were found, the average for each was calculated using `ratingAverage` and saved in the dictionary with movies. In order to prevent `KeyErrors`, if there was a movie with no average for women, it was set to negative infinity. [2] Since we were only interested in the highest ratings, it had no affect on the results. Then the results were printed using `printRatings`, as shown in Listing 1, for the top highest average. Listing 3 is the `loadUserInfo` and `getRatingsByGender` modules.

```

1 def loadUserInfo(path='data'): #load users with info and movie ratings
2     # Load data
3     prefs = {}
4     for line in open(path + '/u.data'):
5         (user, movieid, rating, ts) = line.split('\t')
6         prefs.setdefault(user, {})
7         prefs[user][movieid] = float(rating)
8     for line in open(path + '/u.user'):
9         (user, age, gender, job, zip) = line.split('|')
10        prefs[user]['info'] = {}
11        prefs[user]['info']['age'] = int(age)
12        prefs[user]['info']['gender'] = gender
13    return prefs
14
15 def getRatingsByGender(n=5): #get top average, most rated, ratings by gender
16     rate = loadMovieRatings()
17     # Get ratings per gender
18     users = loadUserInfo()
19     women = {}
20     men = {}
21     for r in rate:
22         women[r] = {}
23         men[r] = {}
24     for u in users:
25         if users[u]['info']['gender'] == 'F':
26             for r in rate:
27                 if users[u].has_key(r):
28                     if women[r].has_key('ratings'):
29                         women[r]['ratings'].append(users[u][r])
30                     else:

```

```

31         women[r]['ratings'] = []
32         women[r]['ratings'].append(users[u][r])
33     else:
34         for r in rate:
35             if users[u].has_key(r):
36                 if men[r].has_key('ratings'):
37                     men[r]['ratings'].append(users[u][r])
38                 else:
39                     men[r]['ratings'] = []
40                     men[r]['ratings'].append(users[u][r])
41 # Get the average ratings per movie per gender
42 for w in women:
43     if women[w].has_key('ratings'):
44         rate[w]['women'] = ratingAverage(women[w])
45 for m in men:
46     if men[m].has_key('ratings'):
47         rate[m]['men'] = ratingAverage(men[m])
48 for r in rate:
49     if rate[r].has_key('women'):
50         continue
51     else:
52         rate[r]['women'] = 0
53 for r in rate:
54     if rate[r].has_key('men'):
55         continue
56     else:
57         rate[r]['men'] = 0
58 # Print results per gender
59 print '\nTop ' + str(n) + ' movies Rated by Women:'
60 printRatings(rate.items(), 'women', n)
61 print '\nTop ' + str(n) + ' movies Rated by Men:'
62 printRatings(rate.items(), 'men', n)

```

Listing 3: Get Top Ratings by Gender Module

1.3.2 RESULT

Surprisingly, only one movie, Prefontaine (1997), was in both highest average list and the women's top ratings list. This leads me to believe that men and women tend to disagree on movie ratings. There were also more than five, but the output says "Top 5" to meet the question requirement. Figure 1.3 shows the list.

Title	Average Rating
Someone Else's America (1995)	5.0
Foreign Correspondent (1940)	5.0
Visitors, The (Visiteurs, Les) (1993)	5.0
Stripes (1981)	5.0
Telling Lies in America (1997)	5.0
Year of the Horse (1997)	5.0
Faster Pussycat! Kill! Kill! (1965)	5.0
Everest (1998)	5.0
Maya Lin: A Strong Clear Vision (1994)	5.0
Prefontaine (1997)	5.0
Mina Tannenbaum (1994)	5.0

Figure 1.3: Highest Average Ratings by Women

1.4 HIGHEST AVERAGE BY MEN

What 5 movies were rated the highest on average by men? Show the movies and their ratings sorted by ratings.

1.4.1 SOLUTION

The same module was used for men as it was for women. A separate dictionary was created for men, to store the movie ratings and average. If a user's gender was not F, the movie list was checked if the user had rated it and if so, the rating was added to the men dictionary. Once all the men ratings were found, the average for each was calculated using `ratingAverage` and saved in the dictionary with movies. In order to prevent `KeyErrors`, if no man rated a movie, it was set to negative infinity, as with the code for women. [2] The code for men was included in listing 3.

1.4.2 RESULT

I found that there were more top rated movies by men that matched the highest average movies from question 1. This tells me that women must not have rated the top rated movies by men that also appear on the highest average list. Prefontaine (1997) was on the list for both men and women, so everyone rated that movie equally high. Again, there were more than five, but for the purposes of the assignment, the output still states "Top 5." Figure 1.4 shows the output.

Top 5 movies Rated by Men:	
Title	Average Rating
-----	-----
Entertaining Angels: The Dorothy Day Story (1996)	5.0
Letter From Death Row, A (1998)	5.0
Hugo Pool (1997)	5.0
Leading Man, The (1996)	5.0
Quiet Room, The (1996)	5.0
Love Serenade (1996)	5.0
Star Kid (1997)	5.0
Great Day in Harlem, A (1994)	5.0
They Made Me a Criminal (1939)	5.0
Delta of Venus (1994)	5.0
Saint of Fort Washington, The (1993)	5.0
Aiqing wansui (1994)	5.0
Little City (1998)	5.0
Santa with Muscles (1996)	5.0
Prefontaine (1997)	5.0
Marlene Dietrich: Shadow and Light (1996)	5.0

Figure 1.4: Highest Average Ratings by Men

1.5 TOP GUN CORRELATION

What movie received ratings most like Top Gun? Which movie received ratings that were least like Top Gun (negative correlation)?

1.5.1 SOLUTION

In order to find a movie with the closest correlation to Top Gun, `calculateSimilarItems`, which was included in the original code, was used. The original was written only for movies, so it was modified to allow for finding similar users as well. The statement to invert the dictionary was removed and I made sure I was sending the correct dictionary to be used. Additionally, it was set to use the `sim_pearson` function, for correlation. To get the lowest negative correlation, two additional functions were written similarly. The main difference was `botMatches` was written to return the worst correlation. The modules for `sim_pearson` was unchanged. The `loadMovie` was used to retrieve with the user rating list, per movie, which is the reverse of the original `loadMovieLens`. This allowed for the correlation to be calculated based on user ratings per movie. [4] Top Gun had to be found in the movie list, so the movie most similar and dissimilar could be retrieved from the results of `calculateSimilarItems` and

calculateDissimilarItems. Listing 4 is the functions written to get the Top Gun correlation.

```
1 def topMatches(prefs, original, n=5, similarity=sim_pearson):
2     '''
3     Returns the best matches for an item from the prefs dictionary.
4     Number of results and similarity function are optional params.
5     '''
6     scores = [(similarity(prefs, original, other), other) for other in prefs
7                if other != original]
8     scores.sort()
9     scores.reverse()
10    return scores[0:n]
11
12 def botMatches(prefs, original, n=5, similarity=sim_pearson):
13     '''
14     Returns the worst matches for an item from the prefs dictionary.
15     Number of results and similarity function are optional params.
16     '''
17     scores = [(similarity(prefs, original, other), other) for other in prefs
18                if other != original]
19     scores.sort()
20    return scores[0:n]
21
22 def calculateSimilarItems(prefs, n=10):
23     '''
24     Create a dictionary of items showing which other items they are
25     most similar to.
26     '''
27     result = {}
28     c = 0
29     for item in prefs:
30         # Status updates for large datasets
31         c += 1
32         if c % 100 == 0:
33             print '%d / %d' % (c, len(prefs))
34         # Find the most similar items to this one
35         scores = topMatches(prefs, item, n=n, similarity=sim_distance)
36         result[item] = scores
37     return result
38
39 def calculateDissimilarItems(prefs, n=10):
40     '''
41     Create a dictionary of items showing which other items they are
42     most similar to.
43     '''
44     result = {}
45     c = 0
46     for item in prefs:
47         # Status updates for large datasets
48         c += 1
49         if c % 100 == 0:
50             print '%d / %d' % (c, len(prefs))
51         # Find the most similar items to this one
52         scores = botMatches(prefs, item, n, similarity=sim_distance)
53         result[item] = scores
54     return result
55
56 def loadMovie(path='data'): #movies, with per user ratings
57     # Get movie titles
58     movies = getMovieTitles()
59     # Load data
60     prefs = {}
61     for line in open(path + '/u.data'):
62         (user, movieid, rating, ts) = line.split('\t')
63         prefs.setdefault(movieid, {})
64         prefs[movieid][user] = float(rating)
65     return prefs
66
67 def getTopGunCorrelation(path='data'): #get movies most alike and unlike Top Gun
68     #load list of movies, with per user ratings
69     mList = loadMovie()
70     topGun = 0
```

```

71 |     # Get movie titles
72 |     movies = getMovieTitles()
73 |     for m in movies:
74 |         if 'Top Gun' in movies[m]:
75 |             topGun = m
76 |     sim = calculateSimilarItems(mList, 1)
77 |     neg = calculateDissimilarItems(mList, 1)
78 |     print '\nMovie most like Top Gun: '
79 |     print movies[sim[m][0][1]]
80 |     print '\nMovie least like Top Gun: '
81 |     print movies[neg[m][0][1]]

```

Listing 4: Top Gun Correlation Module

1.5.2 RESULT

Figure 1.5 shows the movies with ratings most like and most unlike Top Gun. I have not seen either of the movies in the output, so I do not have an opinion about the correlation.

```

Movie most like Top Gun:
Last Time I Committed Suicide, The (1997)

Movie least like Top Gun:
Lightning Jack (1994)

```

Figure 1.5: Top Gun Correlation

1.6 TOP RATERS

Which 5 raters rated the most films? Show the raters' IDs and the number of films each rated.

1.6.1 SOLUTION

A new module, loadUserRatings, was written to compile a dictionary of users with a list of their movie ratings. Using this, the top raters were found by using the length of the list of movie ratings. Listing 5 shows the loadUserRatings and getTopRaters module.

```

1 | def loadUserRatings(path='data'): # users with per movie ratings
2 |     # Get movie titles
3 |     movies = getMovieTitles()
4 |     # Load data
5 |     prefs = {}
6 |     for line in open(path + '/u.data'):
7 |         (user, movieid, rating, ts) = line.split('\t')
8 |         prefs.setdefault(user, {})
9 |         prefs[user][movieid] = float(rating)
10 |    return prefs
11 |
12 | def getTopRaters(n=5): #get users that rated the most movies
13 |     users = loadUserRatings()
14 |     rsort = sorted(users.items(), key=lambda x: len(x[1]), reverse=True)
15 |     print '\nTop 5 users by most ratings:'
16 |     print string.ljust('User', 12), string.rjust('Number of Ratings', 3)
17 |     print string.ljust('-----', 12), string.rjust('-----', 3)
18 |     for x in range(0, n):
19 |         print string.ljust(rsort[x][0], 12), string.rjust(str(len(rsort[x][1])), 3)

```

Listing 5: Top Raters Module

1.6.2 RESULT

There were no ties, so the top five were printed. Figure 1.6 shows the user ID and number of films each rated.

```
Top 5 users by most ratings:
User          Number of Ratings
-----
405           737
655           685
13            636
450           540
276           518
```

Figure 1.6: Top Raters

1.7 USER AGREEMENT

Which 5 raters most agreed with each other? Show the raters' IDs and Pearson's r, sorted by r.

1.7.1 SOLUTION

My solution used `calculateSimilarItems` to find the user closest to each. For each user, a chain was created by joining the user closest to it. Each linked `userid` was put into a list and the links with `r` and distance values were put together into a dictionary. The list of links were converted to a set, to keep the values unique, because in testing I found some users linked back to each other, creating a loop. At the end of the assignment I realized I could have grabbed more results and added to the chain with the next closest user, but there was not enough time to test it. The chain was searched for the shortest distance along the chain. Listing 1.7.1 is the module.

```
1 def getUserCorrelation(n=5): #get users that are most correlated with each other
2     '''
3     Finds the 5 users that agree the most by finding the users that are closest
4     '''
5     users = loadUserRatings()
6     alike = calculateSimilarItems(users, 1)
7     min = float('inf')
8     minid = 0
9     chain = {}
10    r = {}
11    for a in alike:
12        sum = 0
13        c = alike[a]
14        chain[a] = {}
15        links = []
16        links.append(a)
17        r[a] = [a, alike[a][0][1], sim_pearson(users, str(a), alike[a][0][1]), sim_distance(users, str(a),
18            alike[a][0][1])]
19        sum = sum + sim_distance(users, str(a), alike[a][0][1])
20        for i in range(0, n):
21            add = links[i]
22            links.append(alike[add][0][1])
23            next = alike[alike[add][0][1]][0][1]
24            sum = sum + sim_distance(users, alike[add][0][1], alike[alike[add][0][1]][0][1])
```

```

24         chain[a]['links'] = list(set(links))
25         chain[a]['distance'] = sum
26     for c in chain:
27         if len(chain[c]['links']) >= 5:
28             if chain[c]['distance'] < min:
29                 min = chain[c]['distance']
30                 minid = c
31     chain[minid]['r'] = {}
32     for i in range(0, len(chain[minid]['links'])):
33         chain[minid]['r'][i] = r[chain[minid]['links'][i]]
34     csort = sorted(chain[minid]['r'].items(), key=lambda x: x[1][2], reverse=True)
35     print '5 closest raters with r values and distance: '
36     for c in csort:
37         print 'User ' + c[1][0] + ' to ' + c[1][1] + ', r = ' + str(c[1][2]) + ', distance= ' + str(c
          [1][3])

```

1.7.2 RESULT

I would think that if two users had ratings close to each other, the r value would be positively correlated. My results did not support this, the Pearson value was calculated to be zero. When testing, this intuition seemed to be correct for chains up to four. This leads me to believe that my algorithm for finding the five closest was not correct. Figure 1.7 shows the result.

```

5 closest raters with r values and distance:
User 796 to 341, r = 0, distance= 0.2
User 908 to 88, r = 0, distance= 1.0
User 95 to 88, r = 0, distance= 1.0
User 341 to 908, r = 0, distance= 1.0
User 88 to 95, r = 0, distance= 1.0

```

Figure 1.7: User Agreement

1.8 USER DISAGREEMENT

Which 5 raters most disagreed with each other (negative correlation)? Show the raters' IDs and Pearson's r, sorted by r.

1.8.1 SOLUTION

The code to find the 5 raters who disagreed was the same as the code for agreement, except using `calculateDissimilarItems`. Listing 6 is the additional code in `getUserCorrelation`.

```

1     #Finds the 5 users that disagree the most by finding the users that are farthest away
2     unlike = calculateDissimilarItems(users, 1)
3     max = float('-inf')
4     maxid = 0
5     uchain = {}
6     ur = {}
7     for u in unlike:
8         sum = 0
9         c = unlike[u]
10        uchain[u] = {}
11        links = []
12        links.append(u)
13        ur[u] = [a, unlike[u][0][1], sim_pearson(users, str(u), unlike[u][0][1]), sim_distance(users, str(u),
          unlike[u][0][1])]
14        sum = sum + sim_distance(users, str(u), unlike[u][0][1])
15        for i in range(0, n):
16            add = links[i]

```

```

17         links.append(unlike[add][0][1])
18         next = unlike[unlike[add][0][1]][0][1]
19         sum = sum + sim_distance(users, unlike[add][0][1], unlike[unlike[add][0][1]][0][1])
20         uchain[u]['links'] = list(set(links))
21         uchain[u]['distance'] = sum
22     for c in uchain:
23         if len(uchain[c]['links']) == 5:
24             if uchain[c]['distance'] > max:
25                 max = uchain[c]['distance']
26                 maxid = c
27     uchain[maxid]['ur'] = {}
28     for i in range(0, len(uchain[maxid]['links'])):
29         uchain[maxid]['ur'][i] = ur[uchain[maxid]['links'][i]]
30     csort = sorted(uchain[maxid]['ur'].items(), key=lambda x: x[1][2], reverse=True)
31     print '5 farthest raters, with r values and distance: '
32     for c in csort:
33         print 'User ' + c[1][0] + ' to ' + c[1][1] + ', r = ' + str(c[1][2]) + ', distance= ' + str(c
            [1][3])

```

Listing 6: User Disagreement

1.8.2 RESULT

My results for disagreement also do not follow what I would expect. All of the Pearson r values were zero, whereas I would expect them to be negative. Figure ?? shows the result.

```

5 farthest raters, with r values and distance:
User 479 to 445, r = 0.15494918816, distance= 0.00636942675159
User 479 to 124, r = 0, distance= 0
User 479 to 100, r = 0, distance= 0
User 479 to 100, r = 0, distance= 0
User 479 to 565, r = 0, distance= 0

```

Figure 1.8: User Disagreement

1.9 MEN BY AGE

What movie was rated highest on average by men over 40? By men under 40?

1.9.1 SOLUTION

This solution was similar to the gender based solution. The modules used to import the data were loadMovieRatings from listing 1 and loadUserInfo from listing 3. Dictionaries were created for both gender based age groups, so the averages could be calculated independently. Then for each user, gender and age were checked, and the rating for each movie was added. Next the averages were calculated using ratingAverage, as in listing 1. If the movie was not rated by that group, it was set to negative infinity, since it would not affect the highest average outcome. Then the results were printed using printRatings, also from listing 1. Listing 7 is the function getRatingsByAge.

```

1 def getRatingsByAge(n=5): #get top average, most rated, ratings by gender
2     rate = loadMovieRatings()
3     users = loadUserInfo()
4     women = {}
5     w40 = {}
6     men = {}
7     m40 = {}
8     for r in rate:
9         women[r] = {}
10        men[r] = {}
11        w40[r] = {}

```

```

12|     m40[r] = {}
13| for u in users:
14|     if users[u]['info']['gender'] == 'F':
15|         if int(users[u]['info']['age']) < 40:
16|             for r in rate:
17|                 if users[u].has_key(r):
18|                     if women[r].has_key('ratings'):
19|                         women[r]['ratings'].append(users[u][r])
20|                     else:
21|                         women[r]['ratings'] = []
22|                         women[r]['ratings'].append(users[u][r])
23|             else:
24|                 for r in rate:
25|                     if users[u].has_key(r):
26|                         if w40[r].has_key('ratings'):
27|                             w40[r]['ratings'].append(users[u][r])
28|                         else:
29|                             w40[r]['ratings'] = []
30|                             w40[r]['ratings'].append(users[u][r])
31|         else:
32|             if int(users[u]['info']['age']) < 40:
33|                 for r in rate:
34|                     if users[u].has_key(r):
35|                         if men[r].has_key('ratings'):
36|                             men[r]['ratings'].append(users[u][r])
37|                         else:
38|                             men[r]['ratings'] = []
39|                             men[r]['ratings'].append(users[u][r])
40|             else:
41|                 for r in rate:
42|                     if users[u].has_key(r):
43|                         if m40[r].has_key('ratings'):
44|                             m40[r]['ratings'].append(users[u][r])
45|                         else:
46|                             m40[r]['ratings'] = []
47|                             m40[r]['ratings'].append(users[u][r])
48| # Get the average ratings per movie per gender
49| for w in women:
50|     if women[w].has_key('ratings'):
51|         rate[w]['women'] = ratingAverage(women[w])
52|     else:
53|         rate[w]['women'] = float('-inf')
54| for m in men:
55|     if men[m].has_key('ratings'):
56|         rate[m]['men'] = ratingAverage(men[m])
57|     else:
58|         rate[m]['men'] = float('-inf')
59| for w in w40:
60|     if w40[w].has_key('ratings'):
61|         rate[w]['w40'] = ratingAverage(w40[w])
62|     else:
63|         rate[w]['w40'] = float('-inf')
64| for m in m40:
65|     if m40[m].has_key('ratings'):
66|         rate[m]['m40'] = ratingAverage(m40[m])
67|     else:
68|         rate[m]['m40'] = float('-inf')
69| # Print results per gender
70| print '\nTop ' + str(n) + ' movies Rated by Men under 40:'
71| printRatings(rate.items(), 'men', n)
72| print '\nTop ' + str(n) + ' movies Rated by Men over 40:'
73| printRatings(rate.items(), 'm40', n)
74| print '\nTop ' + str(n) + ' movies Rated by Women under 40:'
75| printRatings(rate.items(), 'women', n)
76| print '\nTop ' + str(n) + ' movies Rated by Women over 40:'
77| printRatings(rate.items(), 'w40', n)

```

Listing 7: Ratings by Age

1.9.2 RESULT

Although the requirement was for five, there were more that had a average rating of 5.0. There were a few movies from the top rated men's list that were also on the list for men over 40. Figure 1.9 shows the list. However, most of the movies on the men under 40 list, as shown in Figure 1.10, were also on the top rated men's list. This tells me that most of the movies on the top rated men's list were not rated by men over 40. The movie Prefontaine (1997) was on both lists, so again, it was liked by everyone.

Top 5 movies Rated by Men over 40:	
Title	Average Rating
-----	-----
Leading Man, The (1996)	5.0
Faithful (1996)	5.0
Rendezvous in Paris (Rendez-vous de Paris, Les) (1995)	5.0
Aparajito (1956)	5.0
Star Kid (1997)	5.0
Great Day in Harlem, A (1994)	5.0
They Made Me a Criminal (1939)	5.0
Poison Ivy II (1995)	5.0
Two or Three Things I Know About Her (1966)	5.0
Little Princess, The (1939)	5.0
Late Bloomers (1996)	5.0
Solo (1996)	5.0
Grateful Dead (1995)	5.0
Hearts and Minds (1996)	5.0
Little City (1998)	5.0
Boxing Helena (1993)	5.0
World of Apu, The (Apu Sansar) (1959)	5.0
Spice World (1997)	5.0
Double Happiness (1994)	5.0
Bitter Sugar (Azucar Amargo) (1996)	5.0
Prefontaine (1997)	5.0
Marlene Dietrich: Shadow and Light (1996)	5.0

Figure 1.9: Top Ratings Men over 40

Top 5 movies Rated by Men under 40:	
Title	Average Rating
-----	-----
Entertaining Angels: The Dorothy Day Story (1996)	5.0
Letter From Death Row, A (1998)	5.0
Hugo Pool (1997)	5.0
Leading Man, The (1996)	5.0
Quiet Room, The (1996)	5.0
Love Serenade (1996)	5.0
Star Kid (1997)	5.0
Perfect Candidate, A (1996)	5.0
Delta of Venus (1994)	5.0
Love in the Afternoon (1957)	5.0
Saint of Fort Washington, The (1993)	5.0
Aiqing wansui (1994)	5.0
Crossfire (1947)	5.0
Santa with Muscles (1996)	5.0
Magic Hour, The (1998)	5.0
Angel Baby (1995)	5.0
Maya Lin: A Strong Clear Vision (1994)	5.0
Prefontaine (1997)	5.0

Figure 1.10: Top Ratings Men under 40

1.10

What movie was rated highest on average by women over 40? By women under 40?

1.10.1 SOLUTION

The solution for getting the ratings for women by age group was also in `getRatingsByAge`, shown in listing 7. Again, separate dictionaries were created to consolidate ratings and calculate the averages independently.

1.10.2 RESULT

There were few movies in the women over 40 group that were also in the women's top rated list. This was the first list that did not include the movie *Prefontaine* (1997), so women over 40 must not watch this movie. Figure ?? shows the list. As with the men, most of the movies in the women under 40 group were in the women's top rated list, shown in Figure ?. *Prefontaine* (1997) was also in the list, as it was in both groups of men.

Top 5 movies Rated by Women over 40:	
Title	Average Rating
-----	-----
Nightmare Before Christmas, The (1993)	5.0
Letter From Death Row, A (1998)	5.0
Shall We Dance? (1937)	5.0
Night Flier (1997)	5.0
Pocahontas (1995)	5.0
Wrong Trousers, The (1993)	5.0
Swept from the Sea (1997)	5.0
Great Dictator, The (1940)	5.0
Mrs. Winterbourne (1996)	5.0
Safe (1995)	5.0
Top Hat (1935)	5.0
Foreign Correspondent (1940)	5.0
Ma vie en rose (My Life in Pink) (1997)	5.0
Visitors, The (Visiteurs, Les) (1993)	5.0
Grand Day Out, A (1992)	5.0
Funny Face (1957)	5.0
Bent (1997)	5.0
Shallow Grave (1994)	5.0
Tombstone (1993)	5.0
In the Bleak Midwinter (1995)	5.0
U Turn (1997)	5.0
Angel Baby (1995)	5.0
Best Men (1997)	5.0
Band Wagon, The (1953)	5.0
Mina Tannenbaum (1994)	5.0

Figure 1.11: Top Ratings Women over 40

Top 5 movies Rated by Women under 40:	
Title	Average Rating
-----	-----
Nico Icon (1995)	5.0
Backbeat (1993)	5.0
Umbrellas of Cherbourg, The (Parapluies de Cherbourg, Les) (1964)	5.0
Someone Else's America (1995)	5.0
Don't Be a Menace to South Central While Drinking Your Juice in the Hood (1996)	5.0
Stripes (1981)	5.0
Heaven's Prisoners (1996)	5.0
Telling Lies in America (1997)	5.0
Year of the Horse (1997)	5.0
Faster Pussycat! Kill! Kill! (1965)	5.0
Everest (1998)	5.0
Grace of My Heart (1996)	5.0
Wedding Gift, The (1994)	5.0
Horseman on the Roof, The (Hussard sur le toit, Le) (1995)	5.0
Maya Lin: A Strong Clear Vision (1994)	5.0
Prefontaine (1997)	5.0
Mina Tannenbaum (1994)	5.0

Figure 1.12: Top Ratings Women under 40

1.11 PYTHON CODE

```
1#!/usr/bin/python
2# -*- coding: utf-8 -*-
3from math import sqrt
4import operator
5import string
6
7def sim_distance(prefs, p1, p2):
8    '''
9    Returns a distance-based similarity score for person1 and person2.
10    '''
11    # Get the list of shared_items
12    si = {}
13    for item in prefs[p1]:
14        if item in prefs[p2]:
15            si[item] = 1
16    # If they have no ratings in common, return 0
17    if len(si) == 0:
18        return 0
19    # Add up the squares of all the differences
20    sum_of_squares = sum([pow(prefs[p1][item] - prefs[p2][item], 2) for item in
21                           prefs[p1] if item in prefs[p2]])
22    return 1 / (1 + sum_of_squares)
23
24def sim_pearson(prefs, p1, p2):
25    '''
26    Returns the Pearson correlation coefficient for p1 and p2.
27    '''
28    # Get the list of mutually rated items
29    si = {}
30    for item in prefs[p1]:
31        if item in prefs[p2]:
32            si[item] = 1
33    # If they are no ratings in common, return 0
34    if len(si) == 0:
35        return 0
36    # Sum calculations
37    n = len(si)
38    # Sums of all the preferences
39    sum1 = sum([prefs[p1][it] for it in si])
40    sum2 = sum([prefs[p2][it] for it in si])
41    # Sums of the squares
42    sum1Sq = sum([pow(prefs[p1][it], 2) for it in si])
43    sum2Sq = sum([pow(prefs[p2][it], 2) for it in si])
44    # Sum of the products
45    pSum = sum([prefs[p1][it] * prefs[p2][it] for it in si])
46    # Calculate r (Pearson score)
47    num = pSum - sum1 * sum2 / n
48    den = sqrt((sum1Sq - pow(sum1, 2) / n) * (sum2Sq - pow(sum2, 2) / n))
49    if den == 0:
50        return 0
51    r = num / den
52    return r
53
54def getRecommendations(prefs, person, similarity=sim_pearson):
55    '''
56    Gets recommendations for a person by using a weighted average
57    of every other user's rankings
58    '''
59    totals = {}
60    simSums = {}
61    for other in prefs:
62        # Don't compare me to myself
63        if other == person:
64            continue
65        sim = similarity(prefs, person, other)
66        # Ignore scores of zero or lower
67        if sim <= 0:
68            continue
69    for item in prefs[other]:
70        # Only score movies I haven't seen yet
71        if item not in prefs[person] or prefs[person][item] == 0:
```

```

72         # Similarity * Score
73         totals.setdefault(item, 0)
74         # The final score is calculated by multiplying each item by the
75         # similarity and adding these products together
76         totals[item] += prefs[other][item] * sim
77         # Sum of similarities
78         simSums.setdefault(item, 0)
79         simSums[item] += sim
80     # Create the normalized list
81     rankings = [(total / simSums[item], item) for (item, total) in
82                 totals.items()]
83     # Return the sorted list
84     rankings.sort()
85     rankings.reverse()
86     return rankings
87
88 def transformPrefs(prefs):
89     '''
90     Transform the recommendations into a mapping where persons are described
91     with interest scores for a given title e.g. {title: person} instead of
92     {person: title}.
93     '''
94     result = {}
95     for person in prefs:
96         for item in prefs[person]:
97             result.setdefault(item, {})
98             # Flip item and person
99             result[item][person] = prefs[person][item]
100     return result
101
102 def getRecommendedItems(prefs, itemMatch, user):
103     userRatings = prefs[user]
104     scores = {}
105     totalSim = {}
106     # Loop over items rated by this user
107     for (item, rating) in userRatings.items():
108         # Loop over items similar to this one
109         for (similarity, item2) in itemMatch[item]:
110             # Ignore if this user has already rated this item
111             if item2 in userRatings:
112                 continue
113             # Weighted sum of rating times similarity
114             scores.setdefault(item2, 0)
115             scores[item2] += similarity * rating
116             # Sum of all the similarities
117             totalSim.setdefault(item2, 0)
118             totalSim[item2] += similarity
119     # Divide each total score by total weighting to get an average
120     rankings = [(score / totalSim[item], item) for (item, score) in
121                 scores.items()]
122     # Return the rankings from highest to lowest
123     rankings.sort()
124     rankings.reverse()
125     return rankings
126
127 #-----
128 def getMovieTitles(path='data'):
129     # Get movie titles
130     movies = {}
131     for line in open(path + '/u.item'):
132         (id, title) = line.split('|')[0:2]
133         movies[id] = title
134     return movies
135
136 def loadMovieRatings(path='data'): #load movies with title and ratings list
137     # Get movie titles
138     movies = getMovieTitles()
139     # Get all ratings for each movie
140     rate = {}
141     for line in open(path + '/u.data'):
142         (user, movieid, rating, ts) = line.split('\t')
143         rate.setdefault(movieid, {})
144         rate[movieid]['title'] = movies[movieid]
145         if rate[movieid].has_key('ratings'):

```

```

145         rate[movieid]['ratings'].append(float(rating))
146     else:
147         rate[movieid]['ratings'] = []
148         rate[movieid]['ratings'].append(float(rating))
149     return rate
150
151 def ratingAverage(ratings): #get average of a list of movie ratings
152     sum = 0.0
153     for i in range(0, len(ratings['ratings'])):
154         sum = sum + ratings['ratings'][i]
155     avg = sum / len(ratings['ratings'])
156     return avg
157
158 def getTop(sort, type, n=5):
159     max = sort[0][1][type]
160     top = []
161     top.append(sort[0])
162     for s in range(1, len(sort)):
163         if sort[s][1][type] == max:
164             top.append(sort[s])
165         else:
166             break
167     while len(top) < n:
168         max = sort[len(top)][1][type]
169         for s in range(len(top), len(sort)):
170             if sort[s][1][type] == max:
171                 top.append(sort[s])
172             else:
173                 break
174     return top
175
176 def printRatings(items, type, n): #print list of ratings
177     sort = sorted(items, key=lambda x: x[1][type], reverse=True)
178     col_width = 0
179     top = getTop(sort, type, n)
180     # Get longest title for printed column width
181     for x in range(0, len(top)):
182         if len(top[x][1]['title']) > col_width:
183             col_width = len(top[x][1]['title']) + 2
184     # print the result
185     print string.ljust('Title', col_width), string.rjust('Average Rating', 3)
186     print string.ljust('\n-----', col_width), string.rjust('-----', 3)
187     for x in range(0, len(top)):
188         print string.ljust(top[x][1]['title'], col_width), string.rjust(str(top[x][1][type]), 3)
189
190 def getRatings(n=5): #get top average, most rated
191     rate = loadMovieRatings()
192     # Get the average ratings per movie
193     for r in rate:
194         rate[r]['average'] = ratingAverage(rate[r])
195     print 'Top ' + str(n) + ' movies by Average Rating:'
196     printRatings(rate.items(), 'average', n)
197     # Get the most rated movies
198     rsort = sorted(rate.items(), key=lambda x: len(x[1]['ratings']), reverse=True)
199     print '\nTop ' + str(n) + ' movies by most ratings:'
200     col_width = 0
201     for x in range(0, n):
202         if len(rsort[x][1]['title']) > col_width:
203             col_width = len(rsort[x][1]['title']) + 2
204     print string.ljust('Title', col_width), string.rjust('Number of Ratings', 3)
205     print string.ljust('-----', col_width), string.rjust('-----', 3)
206     for x in range(0, n):
207         print string.ljust(rsort[x][1]['title'], col_width), string.rjust(str(len(rsort[x][1]['ratings'])),
208                                     3)
209
210 #-----
209 def loadUserInfo(path='data'): #load users with info and movie ratings
210     # Load data
211     prefs = {}
212     for line in open(path + '/u.data'):
213         (user, movieid, rating, ts) = line.split('\t')
214         prefs.setdefault(user, {})
215         prefs[user][movieid] = float(rating)
216     for line in open(path + '/u.user'):

```

```

217         (user, age, gender, job, zip) = line.split('|')
218         prefs[user]['info'] = {}
219         prefs[user]['info']['age'] = int(age)
220         prefs[user]['info']['gender'] = gender
221     return prefs
222
223 def getRatingsByGender(n=5): #get top average, most rated, ratings by gender
224     rate = loadMovieRatings()
225     # Get ratings per gender
226     users = loadUserInfo()
227     women = {}
228     men = {}
229     for r in rate:
230         women[r] = {}
231         men[r] = {}
232     for u in users:
233         if users[u]['info']['gender'] == 'F':
234             for r in rate:
235                 if users[u].has_key(r):
236                     if women[r].has_key('ratings'):
237                         women[r]['ratings'].append(users[u][r])
238                     else:
239                         women[r]['ratings'] = []
240                         women[r]['ratings'].append(users[u][r])
241             else:
242                 for r in rate:
243                     if users[u].has_key(r):
244                         if men[r].has_key('ratings'):
245                             men[r]['ratings'].append(users[u][r])
246                         else:
247                             men[r]['ratings'] = []
248                             men[r]['ratings'].append(users[u][r])
249     # Get the average ratings per movie per gender
250     for w in women:
251         if women[w].has_key('ratings'):
252             rate[w]['women'] = ratingAverage(women[w])
253     for m in men:
254         if men[m].has_key('ratings'):
255             rate[m]['men'] = ratingAverage(men[m])
256     for r in rate:
257         if rate[r].has_key('women'):
258             continue
259         else:
260             rate[r]['women'] = 0
261     for r in rate:
262         if rate[r].has_key('men'):
263             continue
264         else:
265             rate[r]['men'] = 0
266     # Print results per gender
267     print '\nTop ' + str(n) + ' movies Rated by Women:'
268     printRatings(rate.items(), 'women', n)
269     print '\nTop ' + str(n) + ' movies Rated by Men:'
270     printRatings(rate.items(), 'men', n)
271
272 #
273 def topMatches(prefs, original, n=5, similarity=sim_pearson):
274     """
275     Returns the best matches for an item from the prefs dictionary.
276     Number of results and similarity function are optional params.
277     """
278     scores = [(similarity(prefs, original, other), other) for other in prefs
279               if other != original]
280     scores.sort()
281     scores.reverse()
282     return scores[0:n]
283
284 def botMatches(prefs, original, n=5, similarity=sim_pearson):
285     """
286     Returns the worst matches for an item from the prefs dictionary.
287     Number of results and similarity function are optional params.
288     """
289     scores = [(similarity(prefs, original, other), other) for other in prefs
290               if other != original]

```

```

290     scores.sort()
291     return scores[0:n]
292
293 def calculateSimilarItems(prefs, n=10):
294     '''
295     Create a dictionary of items showing which other items they are
296     most similar to.
297     '''
298     result = {}
299     c = 0
300     for item in prefs:
301         # Status updates for large datasets
302         c += 1
303         if c % 100 == 0:
304             print '%d / %d' % (c, len(prefs))
305         # Find the most similar items to this one
306         scores = topMatches(prefs, item, n=n, similarity=sim_distance)
307         result[item] = scores
308     return result
309
310 def calculateDissimilarItems(prefs, n=10):
311     '''
312     Create a dictionary of items showing which other items they are
313     most similar to.
314     '''
315     result = {}
316     c = 0
317     for item in prefs:
318         # Status updates for large datasets
319         c += 1
320         if c % 100 == 0:
321             print '%d / %d' % (c, len(prefs))
322         # Find the most similar items to this one
323         scores = botMatches(prefs, item, n, similarity=sim_distance)
324         result[item] = scores
325     return result
326
327 def loadMovie(path='data'): #movies, with per user ratings
328     # Get movie titles
329     movies = getMovieTitles()
330     # Load data
331     prefs = {}
332     for line in open(path + '/u.data'):
333         (user, movieid, rating, ts) = line.split('\t')
334         prefs.setdefault(movieid, {})
335         prefs[movieid][user] = float(rating)
336     return prefs
337
338 def getTopGunCorrelation(path='data'): #get movies most alike and unlike Top Gun
339     #load list of movies, with per user ratings
340     mList = loadMovie()
341     topGun = 0
342     # Get movie titles
343     movies = getMovieTitles()
344     for m in movies:
345         if 'Top Gun' in movies[m]:
346             topGun = m
347     sim = calculateSimilarItems(mList, 1)
348     neg = calculateDissimilarItems(mList, 1)
349     print '\nMovie most like Top Gun: '
350     print movies[sim[m][0][1]]
351     print '\nMovie least like Top Gun: '
352     print movies[neg[m][0][1]]
353
354 def loadUserRatings(path='data'): # users with per movie ratings
355     # Get movie titles
356     movies = getMovieTitles()
357     # Load data
358     prefs = {}
359     for line in open(path + '/u.data'):
360         (user, movieid, rating, ts) = line.split('\t')
361         prefs.setdefault(user, {})
362         prefs[user][movieid] = float(rating)

```

```

363     return prefs
364
365 def getTopRaters(n=5): #get users that rated the most movies
366     users = loadUserRatings()
367     rsort = sorted(users.items(), key=lambda x: len(x[1]), reverse=True)
368     print '\nTop 5 users by most ratings:'
369     print string.ljust('User', 12), string.rjust('Number of Ratings', 3)
370     print string.ljust('-----', 12), string.rjust('-----', 3)
371     for x in range(0, n):
372         print string.ljust(rsort[x][0], 12), string.rjust(str(len(rsort[x][1])), 3)
373
374 #
375 def getUserCorrelation(n=5): #get users that are most correlated with each other
376     '''
377     Finds the 5 users that agree the most by finding the users that are closest
378     '''
379     users = loadUserRatings()
380     alike = calculateSimilarItems(users, 1)
381     min = float('inf')
382     minid = 0
383     chain = {}
384     r = {}
385     for a in alike:
386         sum = 0
387         c = alike[a]
388         chain[a] = {}
389         links = []
390         links.append(a)
391         r[a] = [a, alike[a][0][1], sim_pearson(users, str(a), alike[a][0][1]), sim_distance(users, str(a),
392             alike[a][0][1])]
393         sum = sum + sim_distance(users, str(a), alike[a][0][1])
394         for i in range(0, n):
395             add = links[i]
396             links.append(alike[add][0][1])
397             next = alike[alike[add][0][1]][0][1]
398             sum = sum + sim_distance(users, alike[add][0][1], alike[alike[add][0][1]][0][1])
399             chain[a]['links'] = list(set(links))
400             chain[a]['distance'] = sum
401         for c in chain:
402             if len(chain[c]['links']) >= 5:
403                 if chain[c]['distance'] < min:
404                     min = chain[c]['distance']
405                     minid = c
406         chain[minid]['r'] = {}
407         for i in range(0, len(chain[minid]['links'])):
408             chain[minid]['r'][i] = r[chain[minid]['links'][i]]
409         csort = sorted(chain[minid]['r'].items(), key=lambda x: x[1][2], reverse=True)
410         print '5 closest raters with r values and distance: '
411         for c in csort:
412             print 'User ' + c[1][0] + ' to ' + c[1][1] + ', r = ' + str(c[1][2]) + ', distance = ' + str(c
413                 [1][3])
414
415 #Finds the 5 users that disagree the most by finding the users that are farthest away
416 unlike = calculateDissimilarItems(users, 1)
417 max = float('-inf')
418 maxid = 0
419 uchain = {}
420 ur = {}
421 for u in unlike:
422     sum = 0
423     c = unlike[u]
424     uchain[u] = {}
425     links = []
426     links.append(u)
427     ur[u] = [a, unlike[u][0][1], sim_pearson(users, str(u), unlike[u][0][1]), sim_distance(users, str(u),
428         unlike[u][0][1])]
429     sum = sum + sim_distance(users, str(u), unlike[u][0][1])
430     for i in range(0, n):
431         add = links[i]
432         links.append(unlike[add][0][1])
433         next = unlike[unlike[add][0][1]][0][1]
434         sum = sum + sim_distance(users, unlike[add][0][1], unlike[unlike[add][0][1]][0][1])
435         uchain[u]['links'] = list(set(links))
436         uchain[u]['distance'] = sum
437     for c in uchain:

```

```

443     if len(uchain[c]['links']) == 5:
444         if uchain[c]['distance'] > max:
445             max = uchain[c]['distance']
446             maxid = c
447     uchain[maxid]['ur'] = {}
448     for i in range(0, len(uchain[maxid]['links'])):
449         uchain[maxid]['ur'][i] = ur[uchain[maxid]['links'][i]]
450     csort = sorted(uchain[maxid]['ur'].items(), key=lambda x: x[1][2], reverse=True)
451     print '5 farthest raters, with r values and distance: '
452     for c in csort:
453         print 'User ' + c[1][0] + ' to ' + c[1][1] + ', r = ' + str(c[1][2]) + ', distance = ' + str(c
454             [1][3])
455 #-----
456 def getRatingsByAge(n=5): #get top average, most rated, ratings by gender
457     rate = loadMovieRatings()
458     users = loadUserInfo()
459     women = {}
460     w40 = {}
461     men = {}
462     m40 = {}
463     for r in rate:
464         women[r] = {}
465         men[r] = {}
466         w40[r] = {}
467         m40[r] = {}
468     for u in users:
469         if users[u]['info']['gender'] == 'F':
470             if int(users[u]['info']['age']) < 40:
471                 for r in rate:
472                     if users[u].has_key(r):
473                         if women[r].has_key('ratings'):
474                             women[r]['ratings'].append(users[u][r])
475                         else:
476                             women[r]['ratings'] = []
477                             women[r]['ratings'].append(users[u][r])
478             else:
479                 for r in rate:
480                     if users[u].has_key(r):
481                         if w40[r].has_key('ratings'):
482                             w40[r]['ratings'].append(users[u][r])
483                         else:
484                             w40[r]['ratings'] = []
485                             w40[r]['ratings'].append(users[u][r])
486         else:
487             if int(users[u]['info']['age']) < 40:
488                 for r in rate:
489                     if users[u].has_key(r):
490                         if men[r].has_key('ratings'):
491                             men[r]['ratings'].append(users[u][r])
492                         else:
493                             men[r]['ratings'] = []
494                             men[r]['ratings'].append(users[u][r])
495             else:
496                 for r in rate:
497                     if users[u].has_key(r):
498                         if m40[r].has_key('ratings'):
499                             m40[r]['ratings'].append(users[u][r])
500                         else:
501                             m40[r]['ratings'] = []
502                             m40[r]['ratings'].append(users[u][r])
503     # Get the average ratings per movie per gender
504     for w in women:
505         if women[w].has_key('ratings'):
506             rate[w]['women'] = ratingAverage(women[w])
507         else:
508             rate[w]['women'] = float('-inf')
509     for m in men:
510         if men[m].has_key('ratings'):
511             rate[m]['men'] = ratingAverage(men[m])
512         else:
513             rate[m]['men'] = float('-inf')
514     for w in w40:

```



```

505         if w40[w].has_key('ratings'):
506             rate[w]['w40'] = ratingAverage(w40[w])
507         else:
508             rate[w]['w40'] = float('-inf')
509     for m in m40:
510         if m40[m].has_key('ratings'):
511             rate[m]['m40'] = ratingAverage(m40[m])
512         else:
513             rate[m]['m40'] = float('-inf')
514     # Print results per gender
515     print '\nTop ' + str(n) + ' movies Rated by Men under 40:'
516     printRatings(rate.items(), 'men', n)
517     print '\nTop ' + str(n) + ' movies Rated by Men over 40:'
518     printRatings(rate.items(), 'm40', n)
519     print '\nTop ' + str(n) + ' movies Rated by Women under 40:'
520     printRatings(rate.items(), 'women', n)
521     print '\nTop ' + str(n) + ' movies Rated by Women over 40:'
522     printRatings(rate.items(), 'w40', n)

```

Listing 8: Python Recommendations

REFERENCES

- [1] Alok Singhal. sorting a nested dictionary with lists in python. <http://stackoverflow.com/questions/14719654/sorting-a-nested-dictionary-with-lists-in-python>. Accessed: 2014-11-8.
- [2] David Beazley Brian K. Jones. *Python Cookbook*. O'Reilly Media, 3rd edition, 2013.
- [3] Fancier Output Formatting. <https://docs.python.org/release/1.5.1p1/tut/node50.html>. Accessed: 2014-11-9.
- [4] Toby Segaran. *Programming Collective Intelligence*. O'Reilly Media, 2007.