

Assignment Nine

Sybil Melton

December 3, 2014

CONTENTS

1	Blog Term Matrix	3
1.1	Solution	3
2	Dendrograms	7
2.1	Solution	7
3	K-clustering	10
3.1	Solution	10
4	MDS	10
4.1	Solution	10
5	TFIDF	12
5.1	Solution	12
6	Python Code	15

LISTINGS

1	Retrieve Blogs	3
2	Retrieve Wordcounts	4
3	Create Blog Matrix	5
4	Pages vs Number of Blogs R	7
5	Create Dendrograms	7
6	Create K-clusters	10
7	Create MDS	10
8	TFIDF Calculation	12
9	Complete Feed Parser	15

1 BLOG TERM MATRIX

Create a blog-term matrix. Start by grabbing 100 blogs; include:

- <http://f-measure.blogspot.com/>
- <http://ws-dl.blogspot.com/>

and grab 98 more as per the method shown in class.

Use the blog title as the identifier for each blog (and row of the matrix). Use the terms from every item/title (RSS) or entry/title (Atom) for the columns of the matrix. The values are the frequency of occurrence. Essentially you are replicating the format of the "blogdata.txt" file included with the PCI book code. Limit the number of terms to the most "popular" (i.e., frequent) 500 terms, this is *after* the criteria on p. 32 (slide 7) has been satisfied.

Create a histogram of how many pages each blog has (e.g., 30 blogs with just one page, 27 with two pages, 29 with 3 pages and so on).

1.1 SOLUTION

My first step was to grab a list of blogs, using Blogger.com's "next blog" functionality. The blogs were all written to file by net location. After writing the two required blogs, a loop of 200 iterations ran to get additional blogs, using pyCurl to get the net location. [8] Once each blog was retrieved, it was checked to make sure it didn't already exist in the list with the regular expression module. This is the same function used in Assignment One. [2] Once I had the list, I went through it to check for inappropriate content and to remove any picture, video, or foreign language blogs. My final list of blogs was saved as blogs.txt and is included with this report. Listing 1 is the code used to accomplish the first part of this task.

```
1 def check(url):
2     with open("blogs.txt") as f:
3         found = False
4         for line in f: #iterate over the file one line at a time(memory efficient)
5             if re.search(url, line): #if string found is in current line then keep it
6                 found = True
7     return found
8
9 blogger = 'https://www.blogger.com/next-blog?navBar=true&blogID=953024975153422094'
10 bfile = open('blogs.txt', 'w', 0)
11 bfile.write('f-measure.blogspot.com\n')
12 bfile.write('ws-dl.blogspot.com\n')
13 bfile.close()
14 for i in range(1, 200):
15     buffer = BytesIO()
16     c = pycurl.Curl()
17     c.setopt(c.URL, blogger)
18     c.setopt(c.WRITEDATA, buffer)
19     c.setopt(c.FOLLOWLOCATION, True)
20     c.setopt(c.HTTPHEADER, ['Accept-Language: en'])
21     bfile = open('blogs.txt', 'a', 0)
22     try:
23         c.perform()
24         if c.getinfo(c.RESPONSE_CODE) == 200:
25             o = urlparse(c.getinfo(c.EFFECTIVE_URL))
26             ch = check(o.netloc)
27             if not ch:
28                 bfile.write(o.netloc + '\n')
29                 bfile.close()
30             else:
31                 continue
32     c.close()
```

```

33 except pycurl.error, error:
34     errno, errstr = error
35     print 'An error occurred: ', errstr

```

Listing 1: Retrieve Blogs

My next step was to retrieve the blogs. Each blog in “blogs.txt” was added the feed path */feeds/post-s/default* and encoded with query arguments for Atom, in order to get the raw feed. A feed list was kept, for the stop word hack calculation later on in the program. Each feed was retrieved using the given code to create a `getFeed` function, using `getwordcounts` and `getwords` from the PCI textbook to create the wordcounts dictionary of word counts per blog and `apcount` dictionary of counts per word. [10] For testing, the Web Science/Digital Libraries blog was downloaded to make sure the feed was retrieved and all BeautifulSoup was used to look for the link to the next page, if available. [3] A counter kept track of the number of pages in the blog, which was used for debugging and stored in a dictionary. The pages per blog and number of blogs was written to `blog_pages.txt` in order to create the histogram and is included with this report. [7] The `apcount` dictionary was unsorted, so in order to get the top 500 words, I converted it to a sorted list by word count. A counter was added to the code and stopped the loop if 500 was reached. Listing 2 is the code used to for these tasks.

```

1 def getwordcounts(url):
2     #Returns title and dictionary of word counts for an RSS feed
3     # Parse the feed
4     d = feedparser.parse(url)
5     wc = {}
6     # Loop over all the entries
7     for e in d.entries:
8         if 'summary' in e:
9             summary = e.summary
10        else:
11            summary = e.description
12        # Extract a list of words
13        words = getwords(e.title + ' ' + summary)
14        for word in words:
15            wc.setdefault(word, 0)
16            wc[word] += 1
17    return (d.feed.title, wc)
18
19 def getwords(html):
20     # Remove all the HTML tags
21     txt = re.compile(r'<[^\>]+>').sub('', html)
22     # Split words by all non-alpha characters
23     words = re.compile(r'[^A-Z^a-z^+]').split(txt)
24     # Convert to lowercase
25     return [word.lower() for word in words if word != '']
26
27 def getFeed(feedurl):
28     global wordcounts
29     global apcount
30     try:
31         (title, wc) = getwordcounts(feedurl)
32         if title in wordcounts:
33             for (w, c) in wc.iteritems():
34                 if w in wordcounts[title].iteritems():
35                     wordcounts[title][w] += c
36                 else:
37                     wordcounts[title][w] = c
38         else:
39             wordcounts[title] = wc
40         for (word, count) in wc.items():
41             apcount.setdefault(word, 0)
42             if count > 1:
43                 apcount[word] += 1
44     except:
45         print 'Failed to parse feed %s' % feedurl
46
47 pages = {}
48 feedlist = []

```

```

49|apcount = {}
50|wordcounts = {}
51|wordlist = []
52|scheme = 'http://'
53|path = '/feeds/posts/default'
54|query_arg = {'alt' : 'atom'}
55|udata = urllib.urlencode(query_arg)
56|f = open('blogs.txt', 'r', 0)
57|for line in f:
58|    line = line.strip()
59|    data = urllib.urlencode(query_arg)
60|    full_url = scheme+line+path+"?" + udata
61|    feedlist.append(full_url)
62|    r = requests.get(full_url)
63|    getFeed(full_url)
64|    ddata = r.text
65|    soup = BeautifulSoup(ddata)
66|    next = soup.find('link', rel='next')
67|    count = 1
68|    while next:
69|        n= next.get('href')
70|        getFeed(n)
71|        r = requests.get(n)
72|        data = r.text
73|        soup = BeautifulSoup(data)
74|        count = count + 1
75|        if count % 10 == 0:
76|            print '-- parsing ' + line + ' ' + str(count) + ' pages so far'
77|            next = soup.find('link', rel='next')
78|        print 'Finished parsing ' + line + ' ' + str(count) + ' pages'
79|        if str(count) in pages:
80|            pages[str(count)].append(line)
81|        else:
82|            pages[str(count)] = []
83|            pages[str(count)].append(line)
84|f.close()
85|
86|p = open('blog_pages.txt', 'w', 0)
87|p.write('Pages\tNumber of Blogs\n')
88|for (pg, num) in pages.iteritems():
89|    p.write(pg + '\t' + str(len(num)) + '\n')
90|
91|a = sorted(apcount.items(), key=lambda x: x[1], reverse=True)
92|count = 0
93|for (w, bc) in a:
94|    frac = float(bc) / len(feedlist)
95|    if frac > 0.1 and frac < 0.5:
96|        wordlist.append(w)
97|        count = count + 1
98|        if count >= 500:
99|            break

```

Listing 2: Retrieve Wordcounts

At this point I realized I did not have a the titles or subtitles stored. So I wrote a function, getTitle, in order to accomplish this and were stored as a dictionary. The feed had to be parsed again for both title and subtitle. [5] I found a few still had html markup which caused extra lines and illegal unicode characters which had to be removed. [4], [6] The blog matrix was created and saved as blogdata.txt and is included with this report. The code from the textbook was modified to include the subtitle, or part of the subtitle, as long as the length was less than 100 characters. [11] Listing 3 is the code used to create the blog matrix.

```

1|_illegal_xml_chars_RE = re.compile(u'[\x00-\x08\x0b\xa9\x0c\x0e-\x1F\uD800-\uDFFF\uFFFE\uFFFF]')
2|
3|def remove_tags(text):
4|    cleanr =re.compile('<.*?>')
5|    cleantext = re.sub(cleanr, '', text)
6|    cleantext = re.sub("\n", ' ', cleantext)
7|    cleantext = cleantext.strip()

```

```

8     return _illegal_xml_chars_RE.sub(' ', cleantext)
9
10 def getTitles():
11     titles = {}
12     global scheme, path, udata
13     f = open('blogs.txt', 'r', 0)
14     for line in f:
15         line = line.strip()
16         full_url = scheme+line+path+"?" + udata
17         d = feedparser.parse(full_url)
18         titles[d.feed.title] = remove_tags(d.feed.subtitle)
19     return titles
20
21 titles = {}
22 titles = getTitles()
23 # Blog Term Matrix
24 out = file('blogdata.txt', 'w', 0)
25 out.write('Blog')
26 for word in wordlist:
27     out.write('\t%s' % word)
28 out.write('\n')
29
30 for (blog, wc) in wordcounts.items():
31     if titles[blog] != '':
32         blog = blog + ' - ' + titles[blog]
33         blog = blog.replace(u'\u0144', '\n')
34         blog = _illegal_xml_chars_RE.sub(' ', blog)
35         blog = blog.strip()
36         if len(blog) > 100:
37             blog = blog[:99]
38         print blog
39         out.write(blog)
40         for word in wordlist:
41             if word in wc:
42                 out.write('\t%d' % wc[word])
43             else:
44                 out.write('\t0')
45         out.write('\n')
46 out.close()

```

Listing 3: Create Blog Matrix

Finally, the histogram was created in R, to show the number of blogs and pages. Figure 1.1 show there were more blogs with fewer pages, as there were only 12 blogs with greater than 18 pages. Listing 4 is the R code used to create it.




Figure 1.1: Pages vs Number of Blogs

```
1 bdata = read.table("C:/Users/sybil/Documents/2014 Fall/Web Science/A9/output/blog_pages.txt", header=T, sep
  ="\t")
2
3 newdata <- bdata[order(pages),]
4 pages <- c(newdata$Pages)
5 blogs <- c(newdata$Number.of.Blogs)
6
7 barplot(blogs, names.arg=pages, xlab="Pages", ylab="Number of Blogs")
```

Listing 4: Pages vs Number of Blogs R

2 DENDROGRAMS

Create an ASCII and JPEG dendrogram that clusters (i.e., HAC) the most similar blogs (see slides 12 & 13). Include the JPEG in your report and upload the ascii file to github (it will be too unwieldy for inclusion in the report).

2.1 SOLUTION

First I tried to change the provided function to output to a text file, but it was not successful. So I redirected the standard output temporarily to blog_ascii.txt, which is included with this report. [1] The JPEG dendrogram was successful, as shown in Figure 2.1. The python code used is shown in Listing 5.

```
1 old = sys.stdout
2 sys.stdout = open('blog_ascii.txt', 'w', 0)
3 blognames, words, data=clusters.readfile('blogdata.txt')
4 clust = clusters.hcluster(data)
```

```
5 clusters.printclust(clust, labels=blognames)
6 sys.stdout = old
7 # jpeg dendogram
8 clusters.drawdendrogram(clust, blognames, jpeg='blogclust.jpg')
```

Listing 5: Create Dendrograms

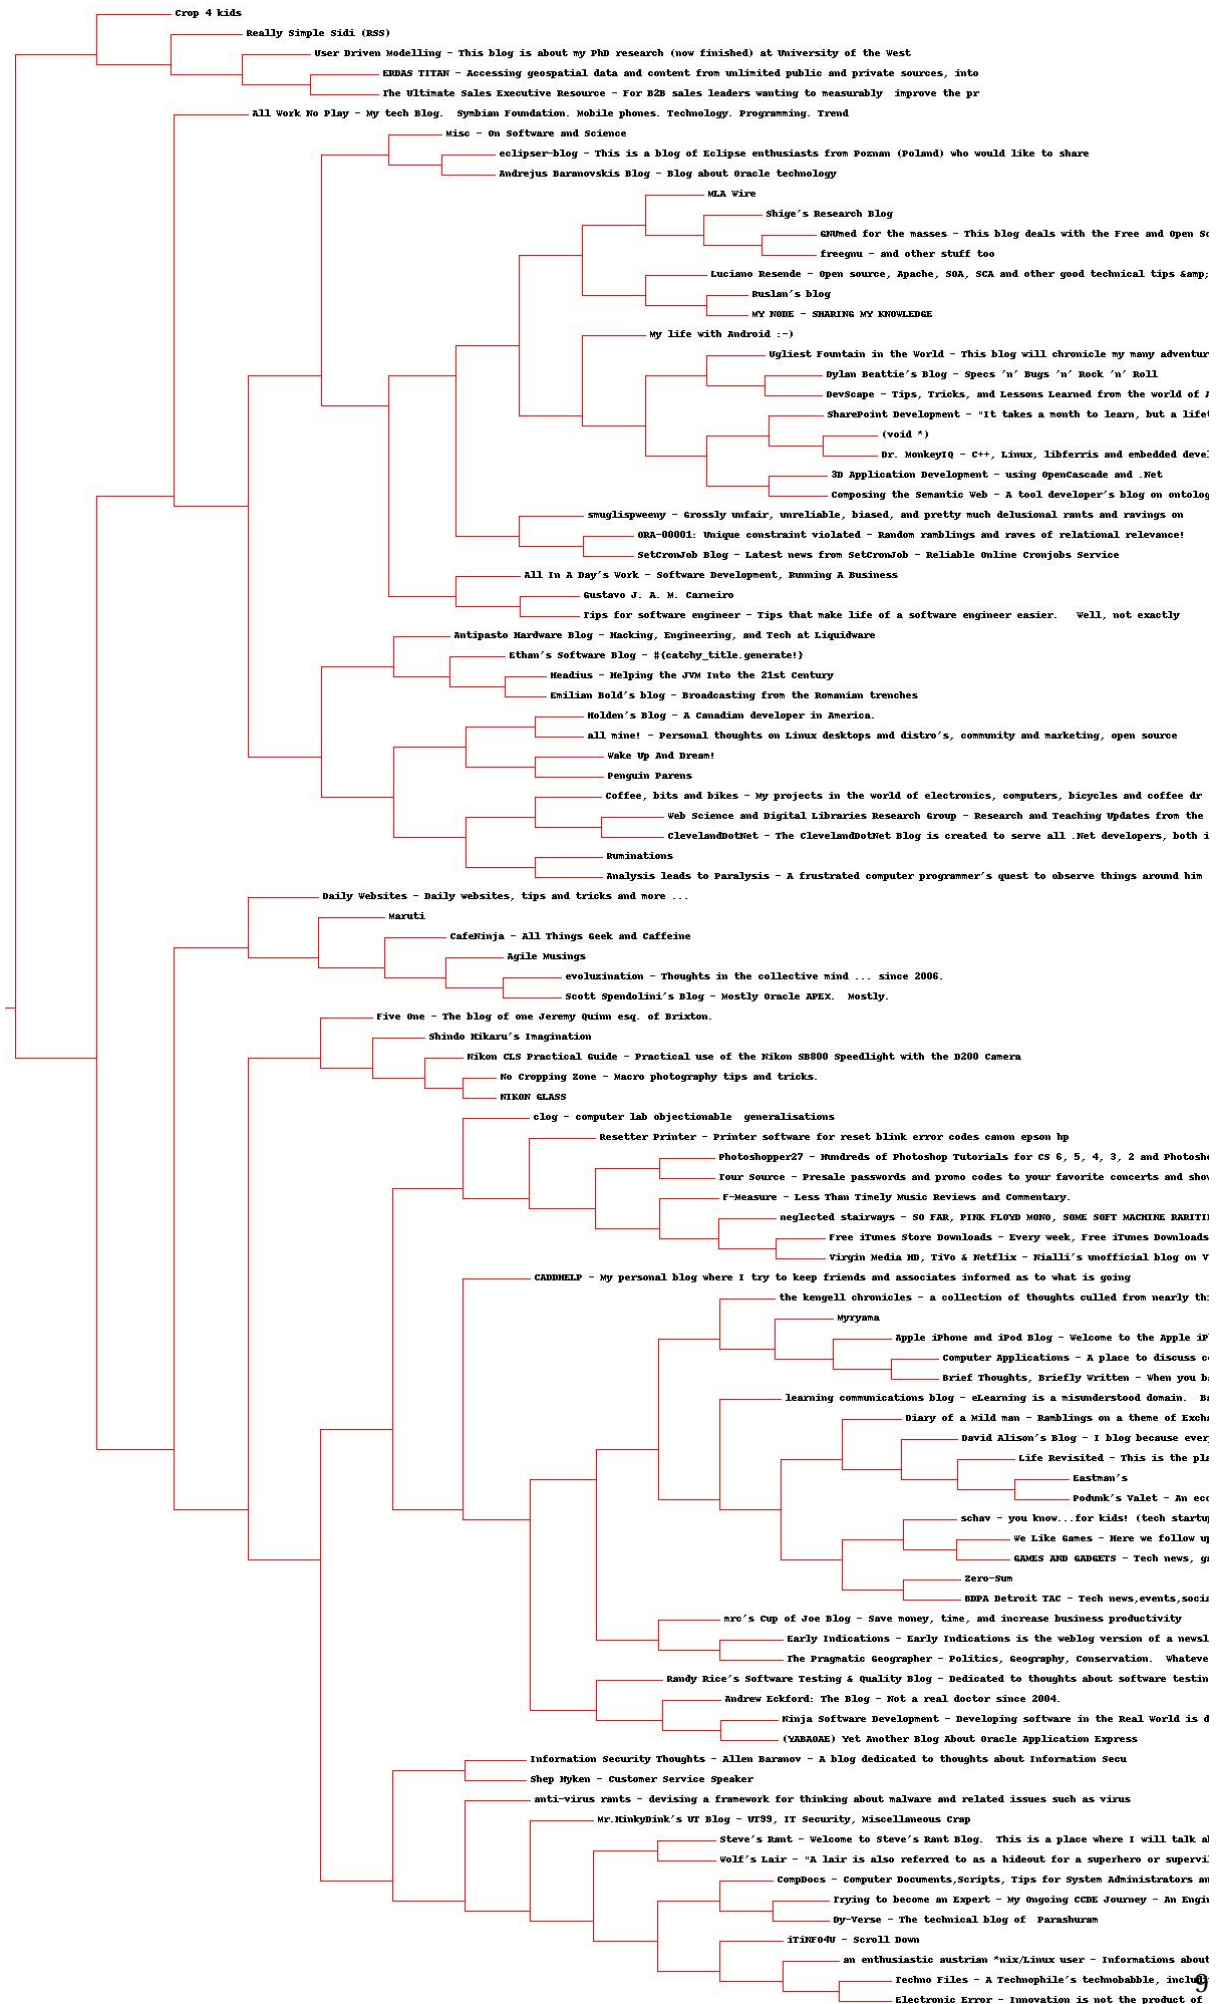


Figure 2.1: JPEG Dendrogram

3 K-CLUSTERING

Cluster the blogs using K-Means, using k=5,10,20. (see slide 18). How many iterations were required for each value of k?

3.1 SOLUTION

To get the k-clustering, I used the provided functions in clusters.py from the PCI textbook, however, I also printed the resulting cluster to a file, blog_k.txt. For K = 5, the code took three iterations, but when K = 10 and 20, it took four iterations. With five clusters, there was a group of 'tips and tricks' blogs, three groups of technology blogs, for gadgets and software development, and a group of hobbies - photoshop, photography, concerts, etc. When K was larger, the technology blogs were split up into smaller clusters, while the other clusters seemed to have stuck together. The python code used to accomplish this is in Listing 6 The screen output from the code was saved as k_run.txt

```
1 def printClusters(kclust, out):
2     for i in range(0,len(kclust)):
3         klist = [blognames[r] for r in kclust[i]]
4         for x in range(0, len(klist)):
5             out.write(klist[x] + '\n')
6         out.write('\n')
7
8 k = open('blog_k.txt', 'w', 0)
9 k.write('K = 5\n')
10 kclust=clusters.kcluster(data,k=5)
11 printClusters(kclust, k)
12 k.write('K = 10\n')
13 kclust=clusters.kcluster(data,k=10)
14 printClusters(kclust, k)
15 k.write('K = 20\n')
16 kclust=clusters.kcluster(data,k=20)
17 printClusters(kclust, k)
18 k.close()
```

Listing 6: Create K-clusters

4 MDS

Use MDS to create a JPEG of the blogs similar to slide 29. How many iterations were required?

4.1 SOLUTION

In order to create the JPEG, the provided code from the PCI textbook was used. The code used is shown in Listing 7 This section took 265 iterations. Although it wasn't required, the iteration numbers were saved as mds.txt, which also made it easier to count the iterations.

```
1 old = sys.stdout
2 sys.stdout = open('blog_ascii.txt', 'w', 0)
3 blognames, words, data=clusters.readfile('blogdata.txt')
4 clust = clusters.hcluster(data)
5 clusters.printclust(clust, labels=blognames)
6 sys.stdout = old
7 # jpeg dendogram
8 clusters.drawdendrogram(clust, blognames, jpeg='blogclust.jpg')
```

Listing 7: Create MDS

CODES AND GADGETS - Tech news, gadget reviews and how-tos on just about anything and everything.

Brief Thoughts, Briefly Written - When you barely have five minutes to yourself, you tend to start daily websites

all mine! - Personal thoughts on various things and a little bit of everything

Apple iPhone and iPod Blog - Welcome to the Apple iPhone and iPod personal review blog! Get news, r

Computer Applications - A place to discuss computer and tech questions.

Haruti

All Work No Play - My tech blog. Synbian Foundation. Mobile phones. Technology. The technoshabbly, including: computer software and hardware, cellphones,

Zero-Sum

Peaback's Valet - An eccentric exploration into the world of a gun-x underachiever.

blog - computer lab objectionable generalisations

We Like Games - Here we follow up with all types of game reviews, game systems

Four Source - Personal

Electronic Error - Innovation is not the product of logical

Eastman's

Web Science and Digital Libraries Research Group - Research and Teaching Updates from the Web Scien

Wolf's Lair - 'A lair is also referred to as a hideout for a superhero or supervillain'

Photoshopper27 - Hundreds of Photoshop Tutorials

Schaefer - you know...for kids!

Guatave J. A. M. Carneiro

Shige's Research Blog

Resetter Printer - Printer software for reset blink error codes canon ep

(void *)

Andrew Eckford: The Blog - Not a real doctor since 2004.

The Pragmatic Geographer - Politics, Geography, Conservation. Whatever

Learning communications blog - Learning is a n

CafeKinja - All Things Geek and Caffeine

Diary of a Wild Man - Ramblings on a theme of Exchange, Parenting, Computers, Cookery and other par

evaluation - Thoughts in the collective mind ... since 2006.

F-Measure - Less Than Timely Music Reviews and Commentary.

ethan's Software Blog - {catchy_title generate!}

Antipasto Hardware Blog - Making, Engineering, and Tech at Liquidware

neglected stairways - 50 FAN, FIVE FLOYD MANS, SOME SHOT MACHINE BARITIES AND THE GOD CAPTAIN BEEF

David Allison's Blog - I blog because everyone is entitled to my opinion

CARROLL - My personal blog where I try to keep friends

Shinde Nikara's Imagination

MLA Wire

Virgin Media HD, FIVE & Netflix - Mialli's unofficial blog on Virgin Media's high definition, system

Spence Scottie's Blog - Specs 'n' Bugs 'n' Back 'n' Ball

Coffee, bits and bikes - My projects in the world of electronics, con

User Driven Modelling -

Free iTunes Store Downloads - Every week, Free iTunes Downloads delivers the free TV, video, music,

Life Revisited - This is the place where you can find

Steve's Rant - Welcome to Steve's Rant Blog. This is a place where I will talk about cool 'geek' t

Really Simple S

Scott Spendolini's Blog - Mostly Oracle APEX. Mostly.

No Cropping Zone - Macro photography tips and tricks.

The Ultimate Sales Executive Resource

Unique constraint violated - Random ramblings and raves of a systems engineer

Upfront Jambain in the World

System's developer instance

linux user - Informations about the

Information Security Thoughts - Allen Ramsey - A blog dedicated to thoughts about Information Secu

Expensive's made with the

Nikon CLS Practical Guide - Practical use of the Nikon SB600 Speed

Hidden's Blog - A Canadian devel

owned for the masses - This blog deals with the Free and Open Source Software owned. It is essent

ServiceScope - Tips, Tricks, and Lessons Learned from the world of Application Development

Smuglywenny - Grossly unfair, unreliable, biased, and pretty much delusional rants and ravings on

Setcrowd Blog - Latest news from Setcrowd - Reliable Online CrowJobs Service

NIKON GLASS

Randy Rice's Software Testing & Quality Blog - Dedicated to thoughts about software testing, an on

Early Indications - Early Indications is the weblog version of a newsletter I've been publishing si

rrv's Cup of Joe Blog - Save money, time, and

Ninja Software Development - Developing software in the Real world is different from all the theory

Agile Mixings

Cleveland@netnet - The Cleveland@netnet Blog is created to serve all .Net developers, both in Clevela

Tips for software engineer - Tips that make life of a softwa

Wake Up And Dream!

(VARIABLE) Yet Another Blog About Oracle Application Express

ROADS TITAN - Accessing geospatial data

Trying to become an Expert - My ongoing OCE Journey - An

Developer's blog on ontology

caliper-blog

my life with Android (-)

Complexes - Computer Documents, Scripts, Tips for System Administrators and Engineers

Microsoft Linux

Shop Nyken - Customer Service Speaks

Eliminations

William's Blog - SHARING MY KNOWLEDGE

Wise - On Software and Science

By-Verse - The technical blog of

William Bold's blog - Broadcasting from the Bonnamim trenches

Luciano Resende

tips amp; tricks on gener

Figure 4.1: MDS

5 TFIDF

Re-run question 2, but this time with proper TFIDF calculations instead of the hack discussed on slide 7 (p. 32). Use the same 500 words, but this time replace their frequency count with TFIDF scores as computed in assignment #3. Document the code, techniques, methods, etc. used to generate these TFIDF values. Upload the new data file to github. Compare and contrast the resulting dendrogram with the dendrogram from question #2.

Note ideally you would not reuse the same 500 terms and instead come up with TFIDF scores for all the terms and then choose the top 500 from that list, but I'm trying to limit the amount of work necessary.

5.1 SOLUTION

TF is calculated as the word frequency divided by the total words in the document. The IDF for each term was calculated as:

$$IDF(\text{term}) = \log_2(\text{total docs in corpus} / \text{docs with term})$$

The total documents in corpus was used as the same as in Assignment Three, which was 42 billion. In order to get the total documents with the term, the python module requests was used to perform the Google search. Beautiful soup found the resultStats, which is where the number of results is stored. My first run was caught by Google's bot detection, so I added a random number sleep element. This took the program a little longer to run, but I was able to get all of the data. The IDF for each term was saved in a dictionary keyed on the term, in order to calculate the TFIDF.

Another blog matrix was created and saved as blog_tfidf.txt. For each word, the TF was calculated using the count from the wordcount dictionary for the term, which was by blog, multiplied by the length of the wordcount entry for the blog, which was the total number of words in the blog. Then TFIDF could be calculated as TF multiplied by IDF, and was rounded to three decimal points. [9]

The last step was to create the new dendrograms. It was similar to the previous dendrogram, if the blogs were not exactly next to each other after the TFIDF calculation, they were in close proximity. This tells me that the "hack" provided by the book was a pretty good estimation for the TFIDF for each blog. Figure 5.1 is the new JPEG dendrogram. Listing 8 is the python code used to accomplish this task.

```
1 google = 'http://www.google.com/search'
2 corp = 42000000000
3 idf = {}
4 for word in wordlist:
5     print 'Searching for ' + word
6     query_arg = {'q' : word}
7     sdata = urllib.urlencode(query_arg)
8     full_url = google+'?' + sdata
9     r = requests.get(full_url)
10    ddata = r.text
11    soup = BeautifulSoup(ddata)
12    results = soup.find('div', id='resultStats')
13    res = (results.text).rsplit()
14    term = float(res[1].replace(',',''))
15    idf[word] = log((corp/term),2)
16    sl = random.randint(1, 100)
17    time.sleep(sl)
18
19 ti = file('blog_tfidf.txt', 'w', 0)
20 ti.write('Blog')
21 for word in wordlist:
22     ti.write('\t%s' % word)
23
24 ti.write('\n')
25
26 for (blog, wc) in wordcounts.items():
```

```

27 | if titles[blog] != '':
28 |     blog = blog + ' - ' + titles[blog]
29 |     blog = blog.replace(u'\u0144', 'n')
30 |     blog = _illegal_xml_chars_RE.sub(' ', blog)
31 |     blog = blog.strip()
32 |     if len(blog) > 100:
33 |         blog = blog[:99]
34 |     ti.write(blog)
35 |     if len(blog) > 100:
36 |         blog = blog[0:99]
37 |     for word in wordlist:
38 |         if word in wc:
39 |             tf = float(wc[word]) / len(wc)
40 |             tfidf = round(tf * idf[word], 3)
41 |             ti.write('\t' + str(tfidf))
42 |         else:
43 |             ti.write('\t0')
44 |     ti.write('\n')
45 |
46 | old = sys.stdout
47 | sys.stdout = open('blog_tfidf_ascii.txt', 'w', 0)
48 | blognames, words, data = clusters.readfile('blog_tfidf.txt')
49 | clust = clusters.hcluster(data)
50 | clusters.printclust(clust, labels=blognames)
51 | sys.stdout = old
52 | clusters.drawdendrogram(clust, blognames, jpeg='tfidfclust.jpg')

```

Listing 8: TFIDF Calculation

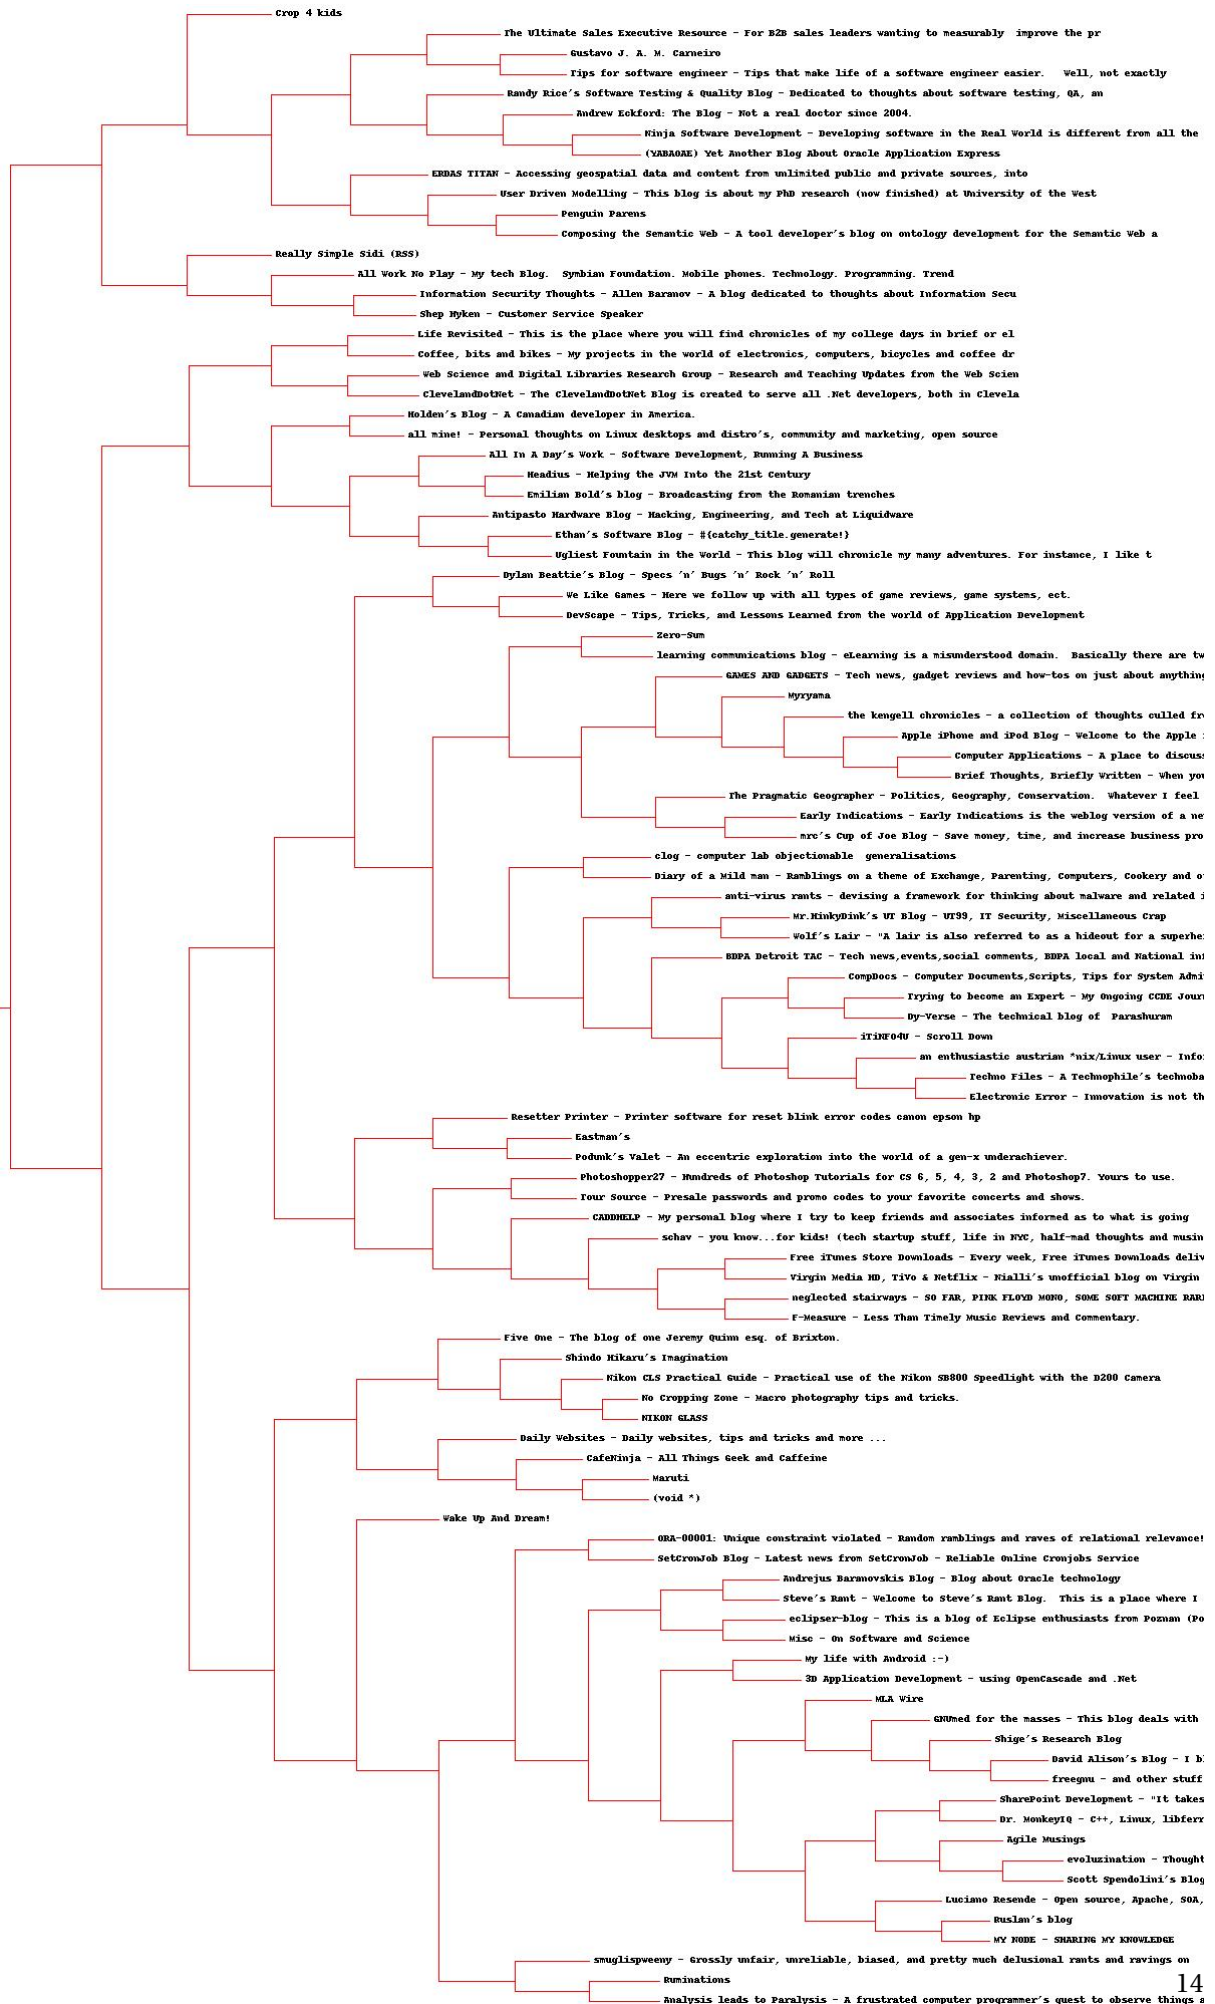


Figure 5.1: TFIDF JPEG Dendrogram

6 PYTHON CODE

```
1 import requests
2 import pycurl
3 import re
4 from urlparse import urlparse
5 import urllib
6 import feedparser
7 from bs4 import BeautifulSoup
8 import clusters
9 import time
10 from math import log
11 import HTMLParser
12 try:
13     from io import BytesIO
14 except ImportError:
15     from StringIO import StringIO as BytesIO
16
17 def check(url):
18     with open("blogs.txt") as f:
19         found = False
20         for line in f: #iterate over the file one line at a time(memory efficient)
21             if re.search(url, line): #if string found is in current line then keep it
22                 found = True
23     return found
24
25 blogger = 'https://www.blogger.com/next-blog?navBar=true&blogID=953024975153422094'
26 bfile = open('blogs.txt', 'w', 0)
27 bfile.write('f-measure.blogspot.com\n')
28 bfile.write('ws-dl.blogspot.com\n')
29 bfile.close()
30 for i in range(1, 200):
31     buffer = BytesIO()
32     c = pycurl.Curl()
33     c.setopt(c.URL, blogger)
34     c.setopt(c.WRITEDATA, buffer)
35     c.setopt(c.FOLLOWLOCATION, True)
36     c.setopt(c.HTTPHEADER, ['Accept-Language: en'])
37     bfile = open('blogs.txt', 'a', 0)
38     try:
39         c.perform()
40         if c.getinfo(c.RESPONSE_CODE) == 200:
41             o = urlparse(c.getinfo(c.EFFECTIVE_URL))
42             ch = check(o.netloc)
43             if not ch:
44                 bfile.write(o.netloc + '\n')
45                 bfile.close()
46             else:
47                 continue
48         c.close()
49 except pycurl.error, error:
50     errno, errstr = error
51     print 'An error occurred: ', errstr
52
53
54
55 def getwordcounts(url):
56     #Returns title and dictionary of word counts for an RSS feed
57     # Parse the feed
58     d = feedparser.parse(url)
59     wc = {}
60     # Loop over all the entries
61     for e in d.entries:
62         if 'summary' in e:
63             summary = e.summary
64         else:
65             summary = e.description
66     # Extract a list of words
67     words = getwords(e.title + ' ' + summary)
68     for word in words:
69         wc.setdefault(word, 0)
70         wc[word] += 1
```



```

71 |     return (d.feed.title , wc)
72 |
73 | def getwords(html):
74 |     # Remove all the HTML tags
75 |     txt = re.compile(r'<[>]+>').sub('', html)
76 |     # Split words by all non-alpha characters
77 |     words = re.compile(r'[^A-Z~a-z]+').split(txt)
78 |     # Convert to lowercase
79 |     return [word.lower() for word in words if word != '']
80 |
81 | def getFeed(feedurl):
82 |     global wordcounts
83 |     global apcount
84 |     try:
85 |         (title , wc) = getwordcounts(feedurl)
86 |         if title in wordcounts:
87 |             for (w, c) in wc.iteritems():
88 |                 if w in wordcounts[title].iteritems():
89 |                     wordcounts[title][w] += c
90 |                 else:
91 |                     wordcounts[title][w] = c
92 |             else:
93 |                 wordcounts[title] = wc
94 |             for (word, count) in wc.items():
95 |                 apcount.setdefault(word, 0)
96 |                 if count > 1:
97 |                     apcount[word] += 1
98 |         except:
99 |             print 'Failed to parse feed %s' % feedurl
100 |
101 | pages = {}
102 | feedlist = []
103 | apcount = {}
104 | wordcounts = {}
105 | wordlist = []
106 | scheme = 'http://'
107 | path = '/feeds/posts/default'
108 | query_arg = {'alt' : 'atom'}
109 | udata = urllib.urlencode(query_arg)
110 | f = open('blogs.txt', 'r', 0)
111 | for line in f:
112 |     line = line.strip()
113 |     data = urllib.urlencode(query_arg)
114 |     full_url = scheme+line+path+"?" + udata
115 |     feedlist.append(full_url)
116 |     r = requests.get(full_url)
117 |     getFeed(full_url)
118 |     ddata = r.text
119 |     soup = BeautifulSoup(ddata)
120 |     next = soup.find('link', rel='next')
121 |     count = 1
122 |     while next:
123 |         n= next.get('href')
124 |         getFeed(n)
125 |         r = requests.get(n)
126 |         data = r.text
127 |         soup = BeautifulSoup(data)
128 |         count = count + 1
129 |         if count % 10 == 0:
130 |             print ' -- parsing ' + line + ' ' + str(count) + ' pages so far'
131 |             next = soup.find('link', rel='next')
132 |     print 'Finished parsing ' + line + ' ' + str(count) + ' pages'
133 |     if str(count) in pages:
134 |         pages[str(count)].append(line)
135 |     else:
136 |         pages[str(count)] = []
137 |         pages[str(count)].append(line)
138 | f.close()
139 |
140 | p = open('blog_pages.txt', 'w', 0)
141 | p.write('Pages\tNumber of Blogs\n')
142 | for (pg, num) in pages.iteritems():
143 |     p.write(pg + '\t' + str(len(num)) + '\n')

```



```

144
145 a = sorted(apcount.items(), key=lambda x: x[1], reverse=True)
146 count = 0
147 for (w, bc) in a:
148     frac = float(bc) / len(feedlist)
149     if frac > 0.1 and frac < 0.5:
150         wordlist.append(w)
151         count = count + 1
152         if count >= 500:
153             break
154
155 _illegal_xml_chars_RE = re.compile(u'[\x00-\x08\x0b\xa9\x0c\x0e-\x1f\uD800-\uDFFF\uFFFE\uFFFF]')
156
157 def remove_tags(text):
158     cleanr = re.compile('<.*?>')
159     cleantext = re.sub(cleanr, '', text)
160     cleantext = re.sub("\n", ' ', cleantext)
161     cleantext = cleantext.strip()
162     return _illegal_xml_chars_RE.sub(' ', cleantext)
163
164 def getTitles():
165     titles = {}
166     global scheme, path, udata
167     f = open('blogs.txt', 'r', 0)
168     for line in f:
169         line = line.strip()
170         full_url = scheme+line+path+"?" + udata
171         d = feedparser.parse(full_url)
172         titles[d.feed.title] = remove_tags(d.feed.subtitle)
173     return titles
174
175 titles = {}
176 titles = getTitles()
177 # Blog Term Matrix
178 out = file('blogdata.txt', 'w', 0)
179 out.write('Blog')
180 for word in wordlist:
181     out.write('\t%s' % word)
182 out.write('\n')
183
184 for (blog, wc) in wordcounts.items():
185     if titles[blog] != '':
186         blog = blog + ' - ' + titles[blog]
187         blog = blog.replace(u'\u0144', '\n')
188         blog = _illegal_xml_chars_RE.sub(' ', blog)
189         blog = blog.strip()
190         if len(blog) > 100:
191             blog = blog[:99]
192         print blog
193         out.write(blog)
194         for word in wordlist:
195             if word in wc:
196                 out.write('\t%d' % wc[word])
197             else:
198                 out.write('\t0')
199         out.write('\n')
200 out.close()
201 # Ascii
202 old = sys.stdout
203 sys.stdout = open('blog_ascii.txt', 'w', 0)
204 blognames, words, data = clusters.readfile('blogdata.txt')
205 clust = clusters.hcluster(data)
206 clusters.printclust(clust, labels=blognames)
207 sys.stdout = old
208 # jpeg dendrogram
209 clusters.drawdendrogram(clust, blognames, jpeg='blogclust.jpg')
210 # K-clustering
211 def printClusters(kclust, out):
212     for i in range(0, len(kclust)):
213         klist = [blognames[r] for r in kclust[i]]
214         for x in range(0, len(klist)):
215             out.write(klist[x] + '\n')
216         out.write('\n')

```

```

217
218 k = open('blog_k.txt', 'w', 0)
219 k.write('K = 5\n')
220 kclust=clusters.kcluster(data,k=5)
221 printClusters(kclust, k)
222 k.write('K = 10\n')
223 kclust=clusters.kcluster(data,k=10)
224 printClusters(kclust, k)
225 k.write('K = 20\n')
226 kclust=clusters.kcluster(data,k=20)
227 printClusters(kclust, k)
228 k.close()
229 # MDS
230 coords = clusters.scaledown(data)
231 clusters.draw2d(coords,blognames,jpeg='blogs2d.jpg')
232 #
233 #Extra credit - TFIDF
234 #
235 google = 'http://www.google.com/search'
236 corp = 42000000000
237 idf = {}
238 for word in wordlist:
239     print 'Searching for ' + word
240     query_arg = {'q' : word}
241     sdata = urllib.urlencode(query_arg)
242     full_url = google+'?' + sdata
243     r = requests.get(full_url)
244     ddata = r.text
245     soup = BeautifulSoup(ddata)
246     results = soup.find('div', id='resultStats')
247     res = (results.text).rsplit()
248     term = float(res[1].replace(',',''))
249     idf[word] = log((corp/term),2)
250     sl = random.randint(1, 100)
251     time.sleep(sl)
252
253 ti = file('blog_tfidf.txt', 'w', 0)
254 ti.write('Blog')
255 for word in wordlist:
256     ti.write('\t%s' % word)
257
258 ti.write('\n')
259
260 for (blog, wc) in wordcounts.items():
261     if titles[blog] != '':
262         blog = blog + ' - ' + titles[blog]
263         blog = blog.replace(u'\u0144', 'n')
264         blog = _illegal_xml_chars_RE.sub(' ', blog)
265         blog = blog.strip()
266         if len(blog) > 100:
267             blog = blog[:99]
268         ti.write(blog)
269         if len(blog) > 100:
270             blog = blog[0:99]
271         for word in wordlist:
272             if word in wc:
273                 tf = float(wc[word]) / len(wc)
274                 tfidf = round(tf * idf[word], 3)
275                 ti.write('\t' + str(tfidf))
276             else:
277                 ti.write('\t0')
278         ti.write('\n')
279
280 old = sys.stdout
281 sys.stdout = open('blog_tfidf_ascii.txt', 'w', 0)
282 blognames, words, data=clusters.readfile('blog_tfidf.txt')
283 clust = clusters.hcluster(data)
284 clusters.printclust(clust, labels=blognames)
285 sys.stdout = old
286 clusters.drawdendrogram(clust, blognames, jpeg='tfidfclust.jpg')

```

Listing 9: Complete Feed Parser

REFERENCES

- [1] Alan Gauld. how to redirect sys.stdout back to screen in IDLE. <https://mail.python.org/pipermail/tutor/2002-November/018792.html>. Accessed: 2014-11-20.
- [2] Ashwini Chaudhary. Python: Checking if string is in a text file. <http://stackoverflow.com/questions/16432203/python-checking-if-string-is-in-a-text-file>. Accessed: 2014-9-20.
- [3] Beautiful soup documentation. <http://www.crummy.com/software/BeautifulSoup/bs4/doc/>. Accessed: 2014-10-03.
- [4] c24b. Python code to remove HTML tags from a string. <http://stackoverflow.com/questions/9662346/python-code-to-remove-html-tags-from-a-string>. Accessed: 2014-11-20.
- [5] Kurt McKee. Common Atom Elements. <https://pythonhosted.org/feedparser/common-atom-elements.html>. Accessed: 2014-11-20.
- [6] Leo Simons. stripping illegal characters out of xml in python. <http://lsimons.wordpress.com/2011/03/17/stripping-illegal-characters-out-of-xml-in-python/>. Accessed: 2014-11-20.
- [7] Philip Southam. 'too many values to unpack', iterating over a dict. key=>string, value=>list. <http://stackoverflow.com/questions/5466618/too-many-values-to-unpack-iterating-over-a-dict-key-string-value-list>. Accessed: 2014-11-20.
- [8] HTTP with Python â&Pound; PycURL by Example. <http://www.angryobjects.com/2011/10/15/http-with-python-pycurl-by-example/>. Accessed: 2014-11-20.
- [9] Python Software Foundation. Built-in Functions. <https://docs.python.org/2/library/functions.html>. Accessed: 2014-11-20.
- [10] Toby Segaran. *Programming Collective Intelligence*. O'Reilly Media, 2007.
- [11] tutorialspoint. Python Strings. http://www.tutorialspoint.com/python/python_strings.htm. Accessed: 2014-11-20.